
MICROSOFT


FABRIC

Bas Land

Your first Lakehouse using Microsoft

DATA:Scotland 2024

Tabular Editor

DATAmasterminds

Waterstons

ADVANCING ANALYTICS

Today

- Core concepts for data warehousing in Fabric

- Your first Fabric Lakehouse:
  - Rulebook
  - ELT process
  - End-to-end solution design with examples

- Takeaways

Core concepts

General design principles

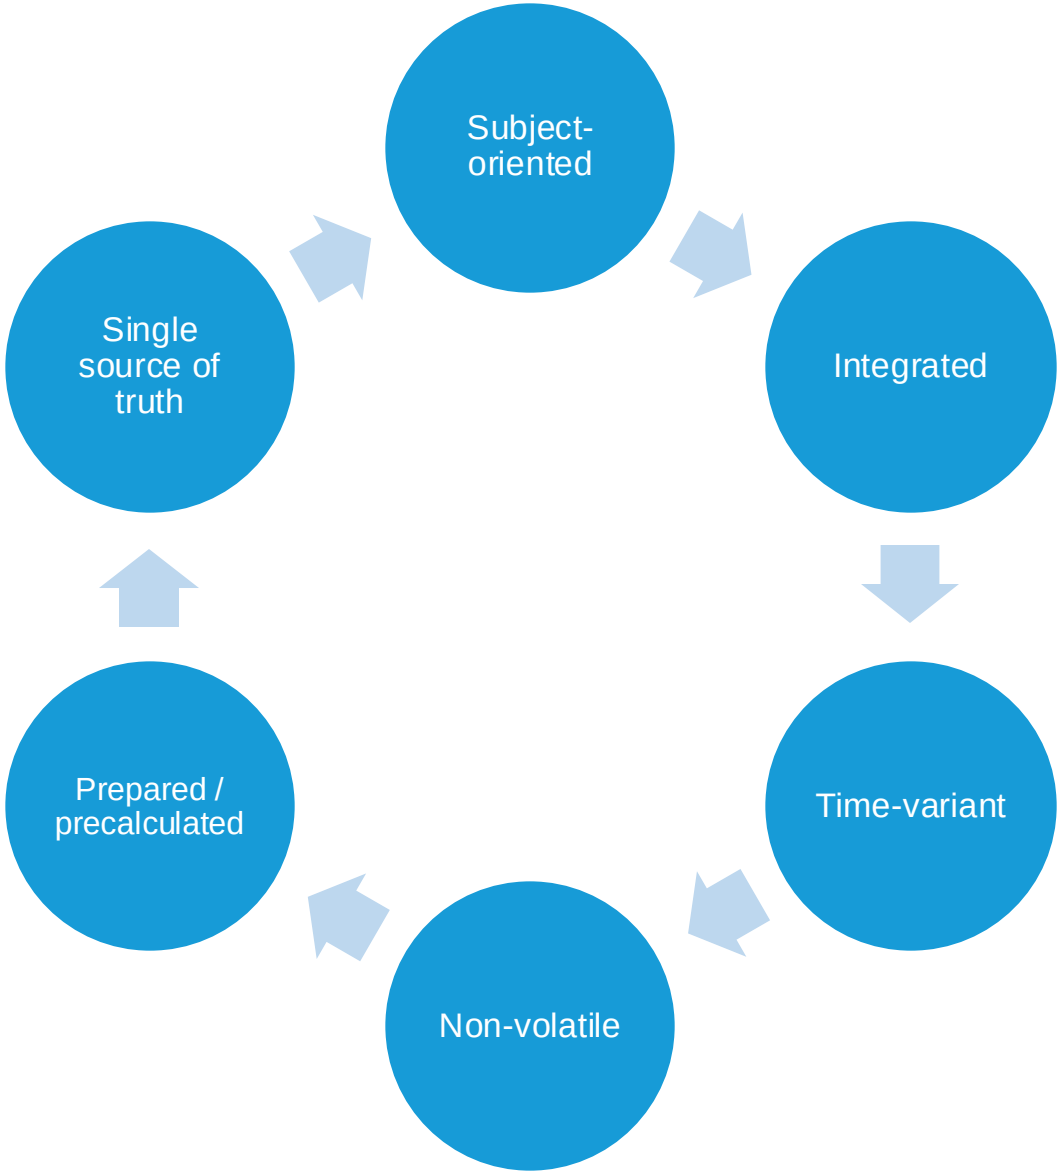

Lakehouse: under the hood

Lakehouse: under the hood

Ingest

Pipelines
Notebooks
Dataflows

Store

OneLake
Data lake blob storage

Process

Notebooks
(PySpark)

Fabric concepts

1. Tenant design

2. Storage

3. Compute

Fabric: tenant design

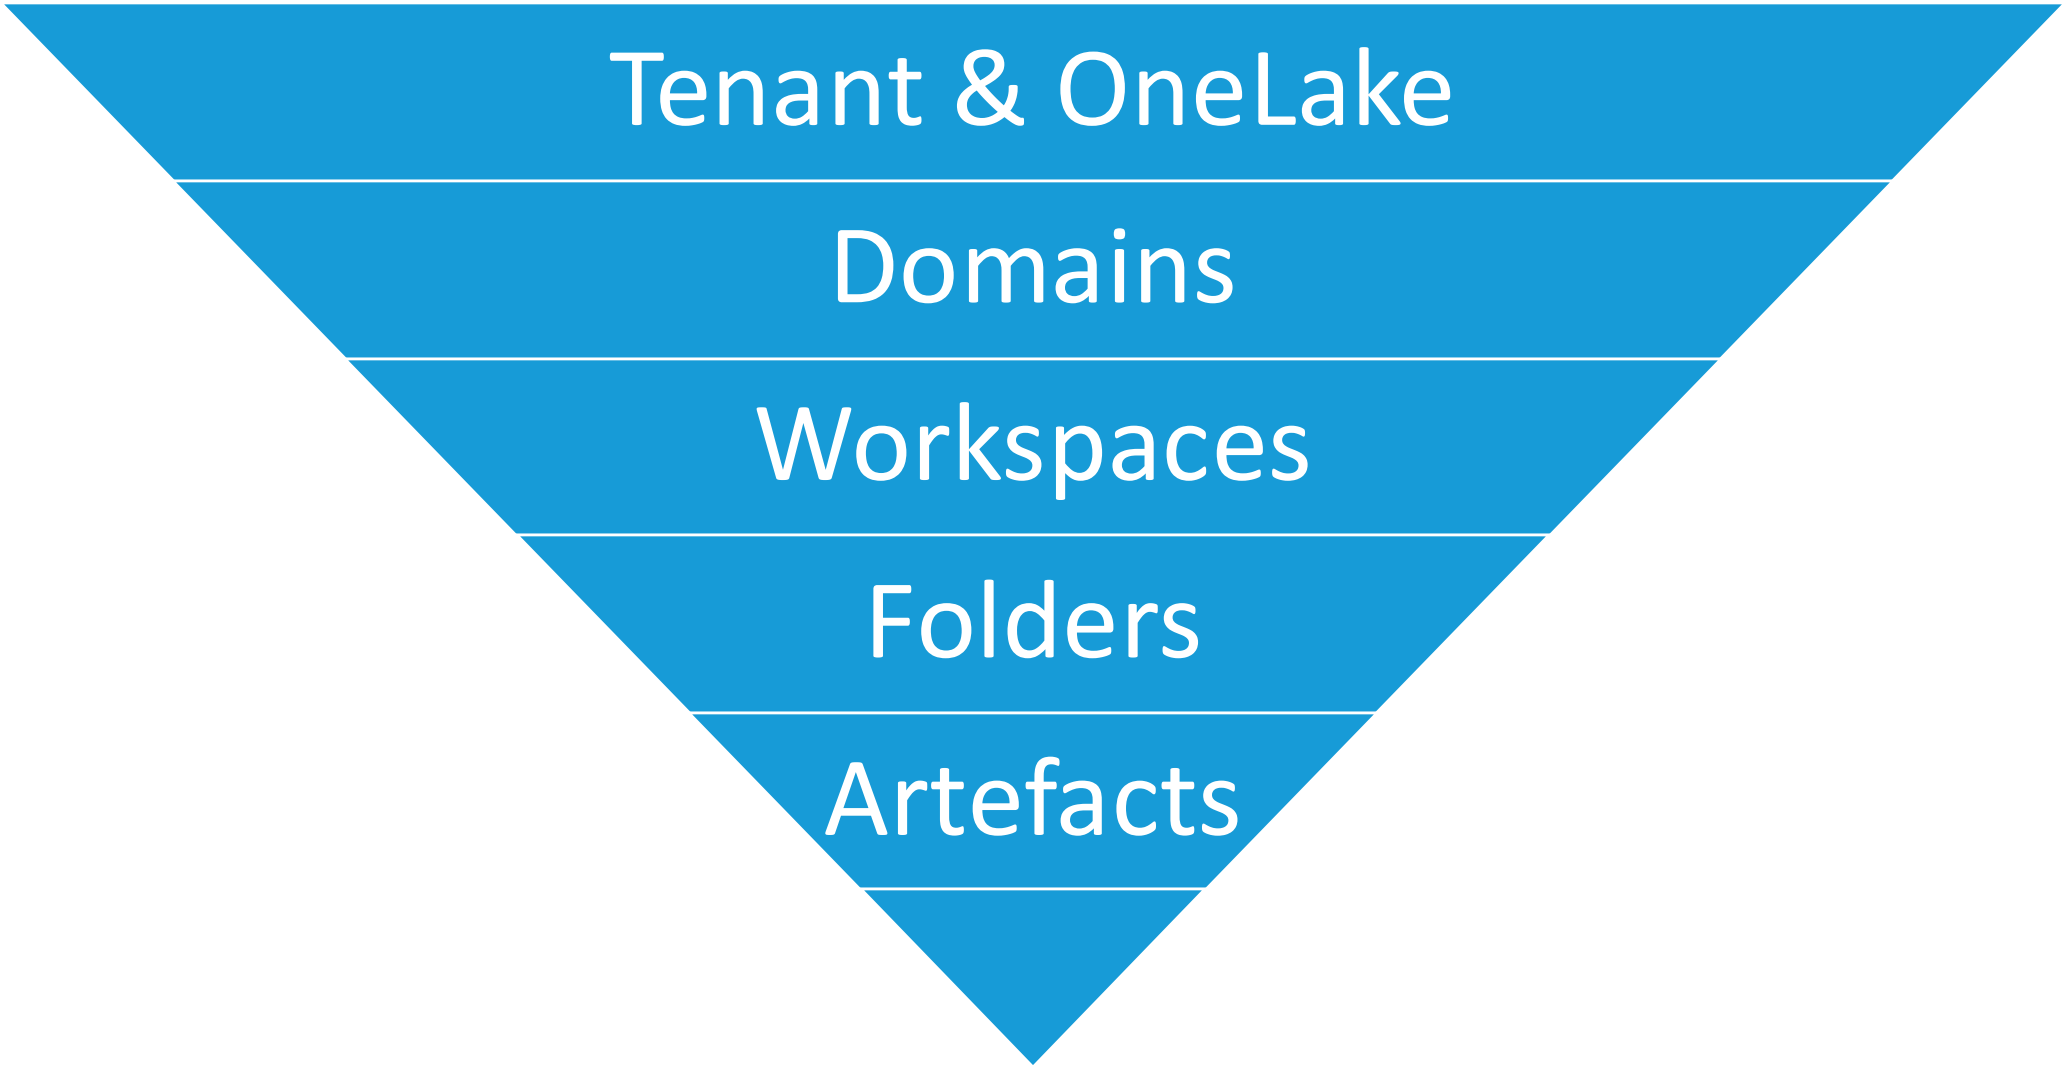

Fabric: tenant design

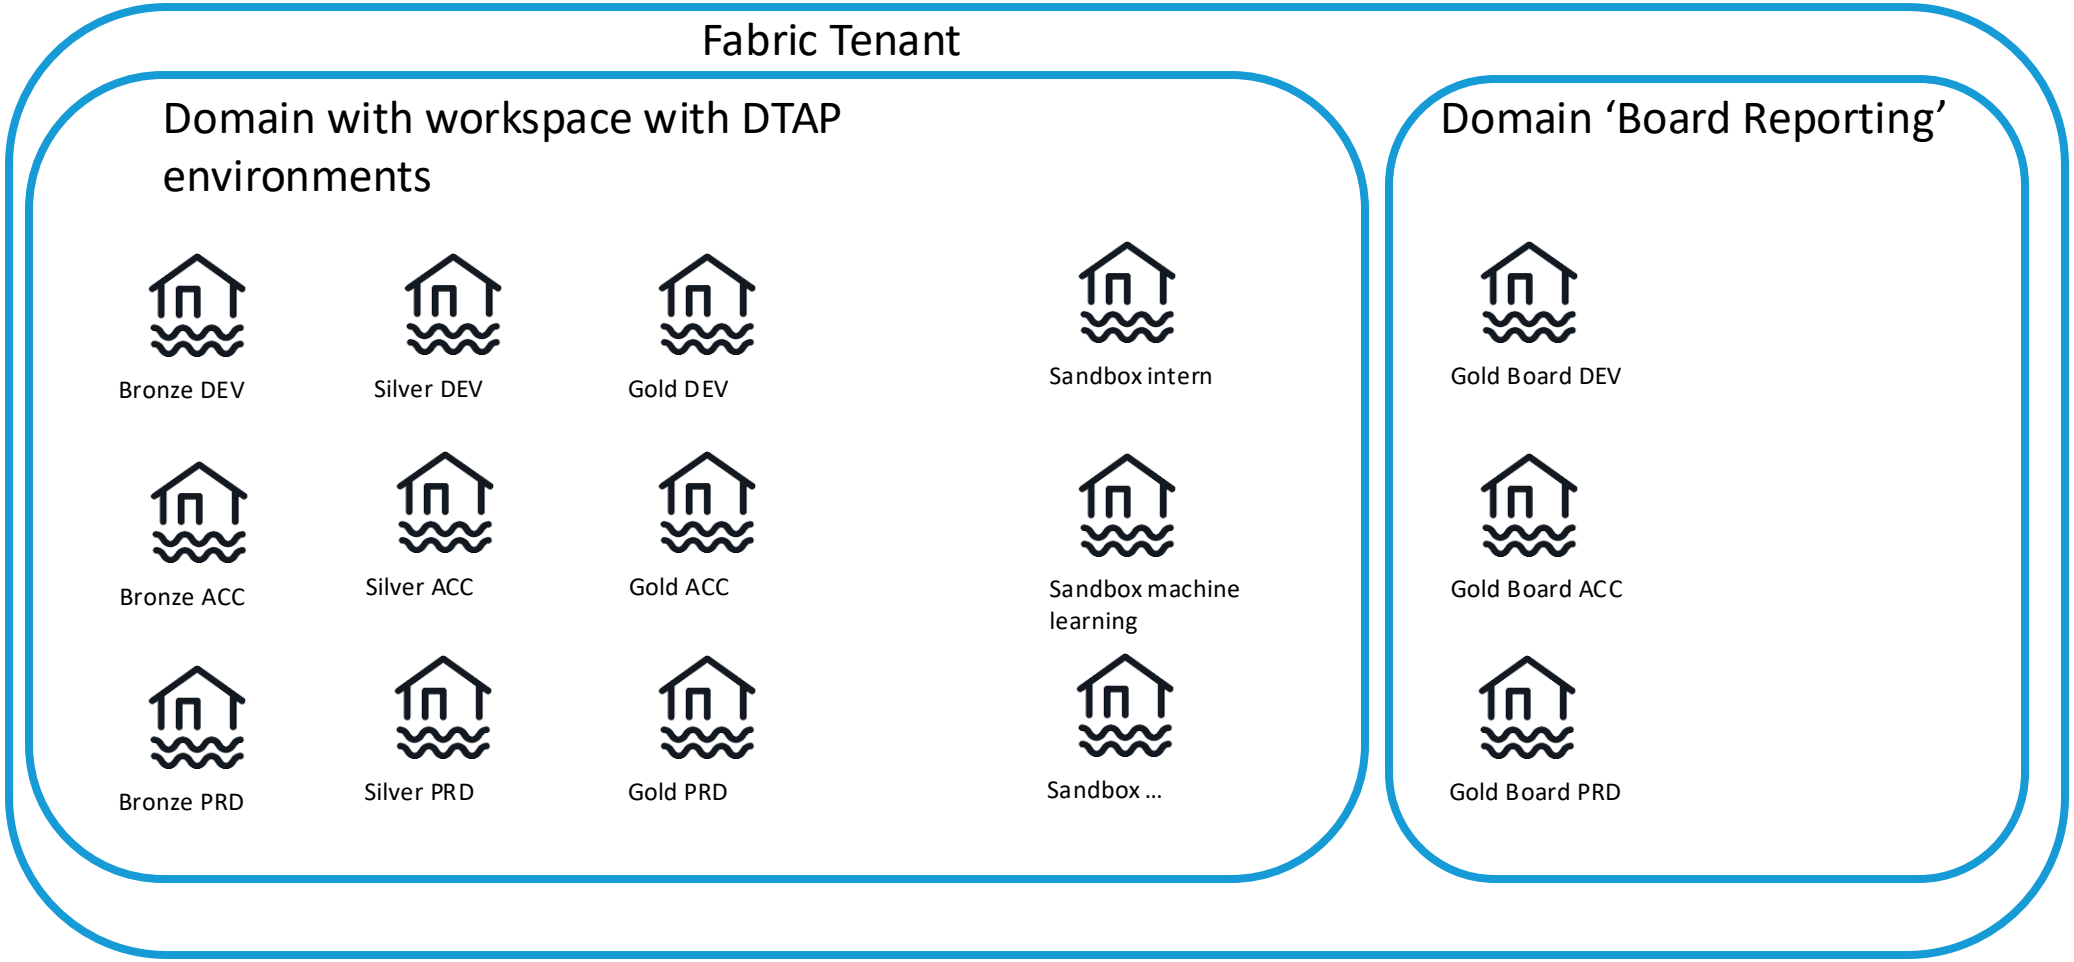

Fabric: storage

- Storage in Fabric is called ‘OneLake’
- It’s data lake storage (blob), with one OneLake per organisation (tenant)
- Storage is cheap! ± 2,2 cent per GB per month

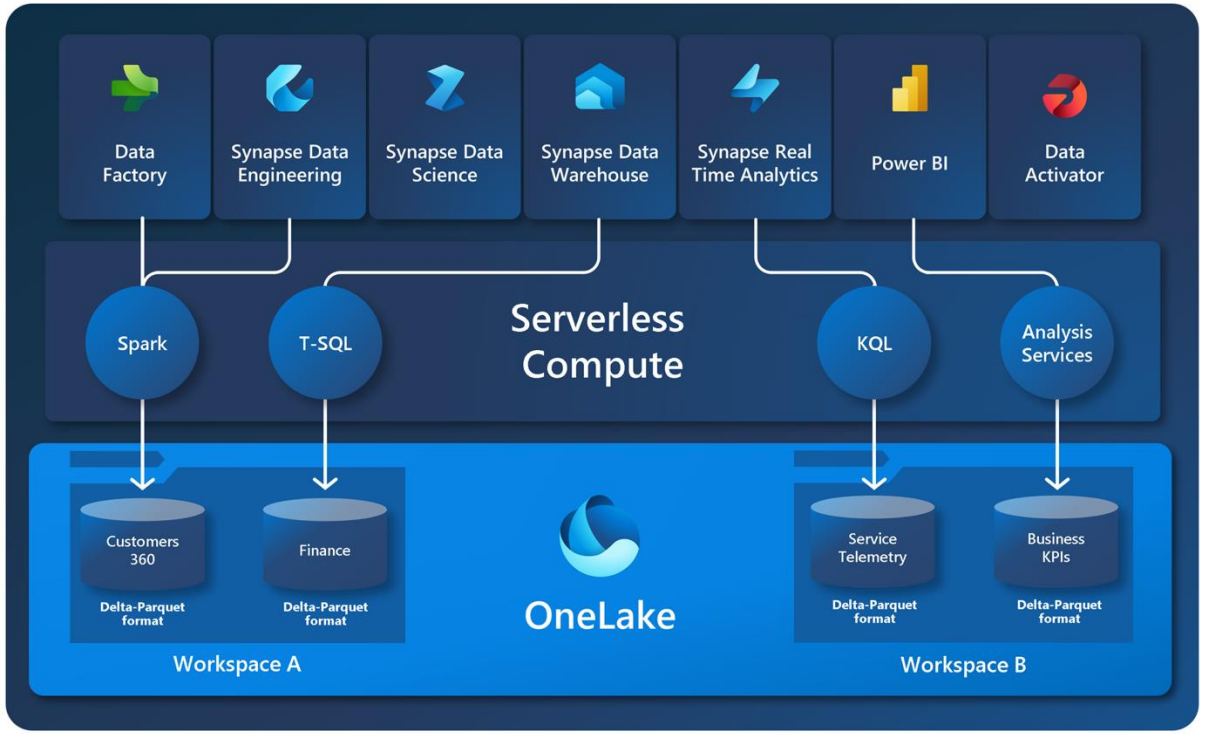

Fabric: compute

- Compute in Fabric is Serverless, and included in your SaaS capacity
- For data engineering purposes, you get a Spark cluster that is preconfigured and can be configured to your custom needs
- Compute = expensive. Be smart about it ☺

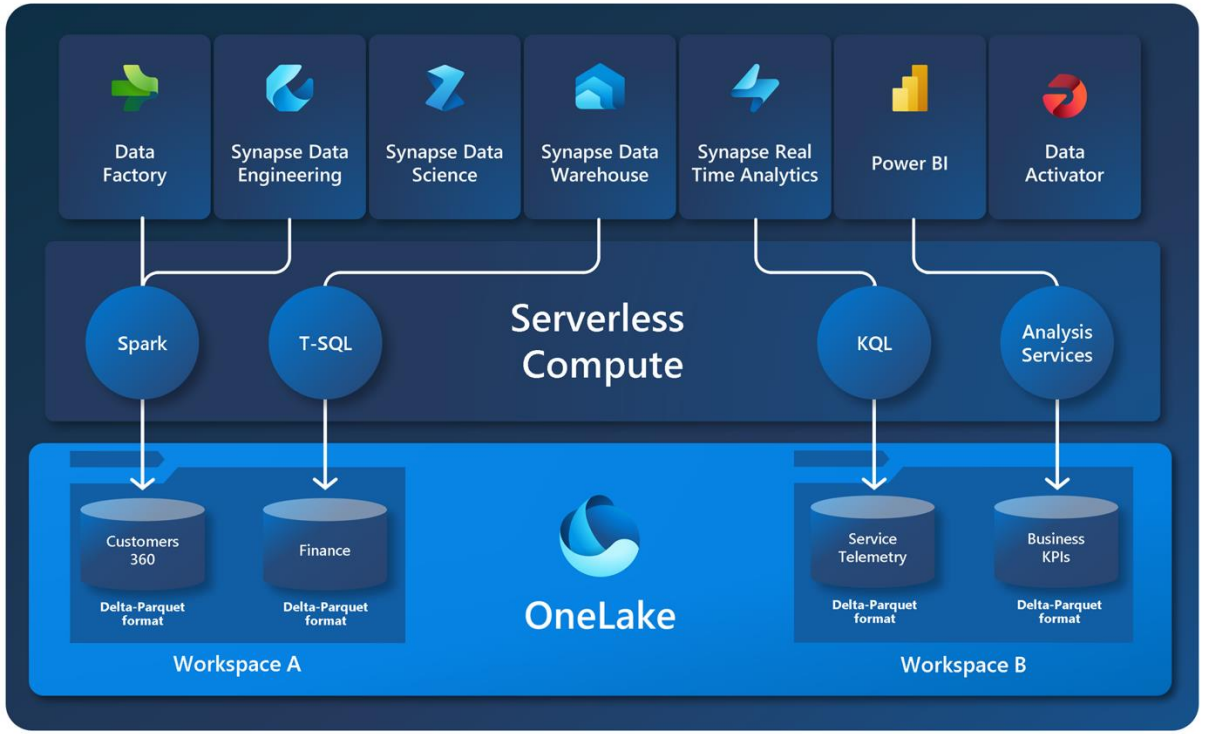

Buying the Fabric compute capacity

- portal.azure.com
- I'd create at least a new Resource Group, maybe even a Subscription (depending on organisation requirements)

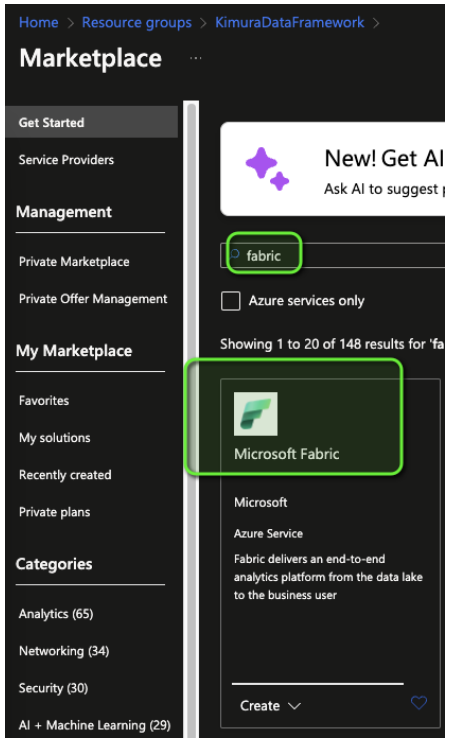

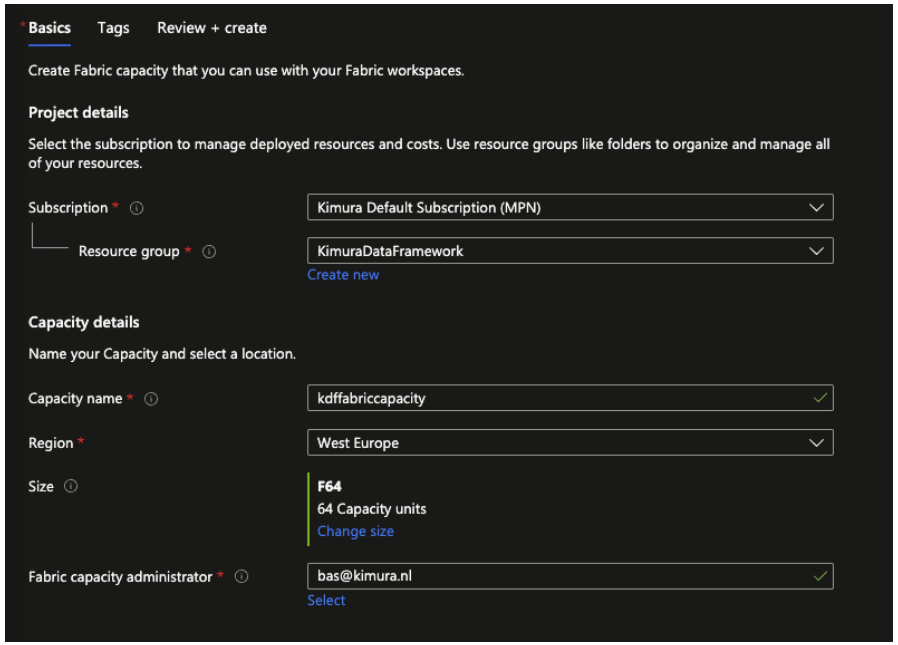

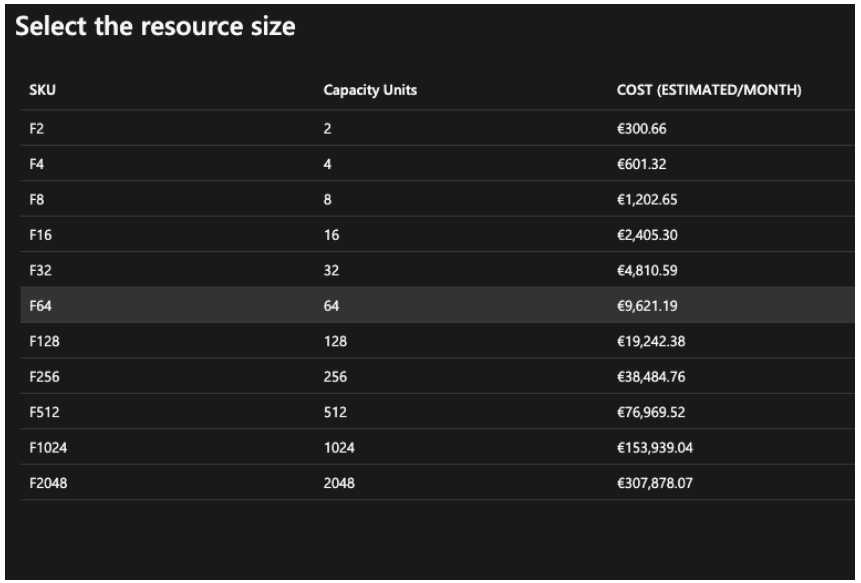

Select the resource size

SKU	Capacity Units	COST (ESTIMATED/MONTH)
F2	2	€300.66
F4	4	€601.32
F8	8	€1,202.65
F16	16	€2,405.30
F32	32	€4,810.59
F64	64	€9,621.19
F128	128	€19,242.38
F256	256	€38,484.76
F512	512	€76,969.52
F1024	1024	€153,939.04
F2048	2048	€307,878.07

Buying the Fabric compute capacity

- After deployment, configure the Capacity Administrator:
- Unfortunately, no Entra ID groups…

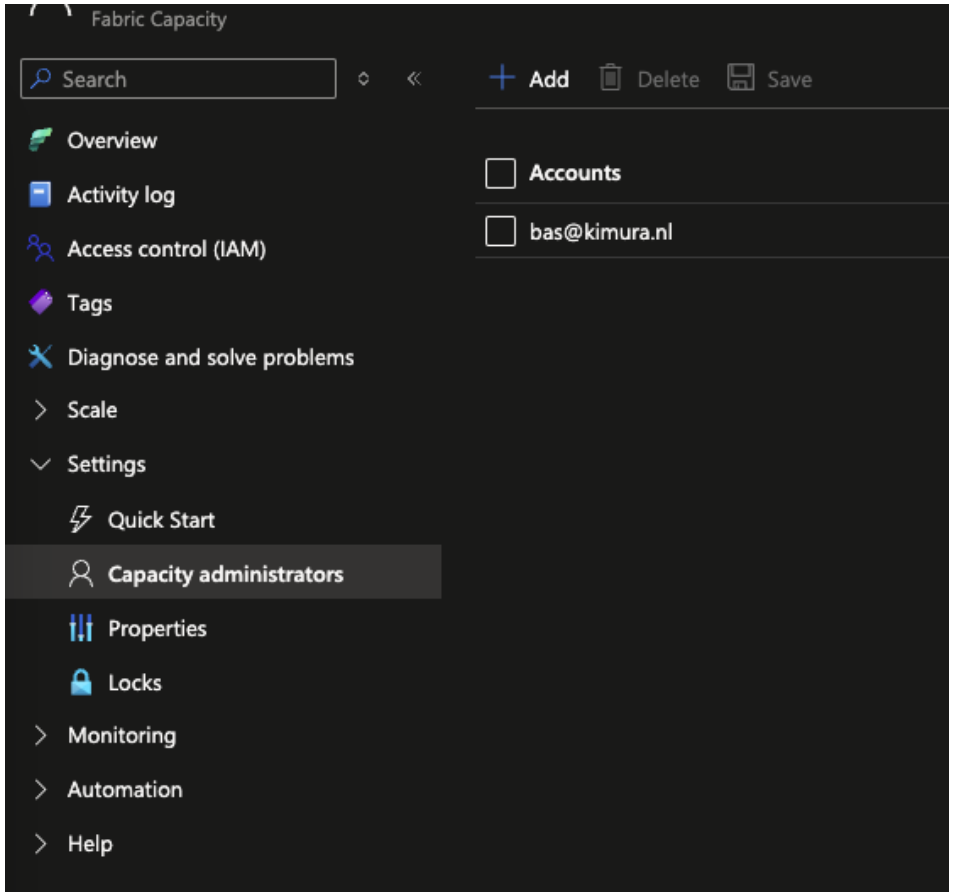

Creating the first Fabric Workspace

- Head over to app.powerbi.com
- Log in with the same account that is the Capacity Administrator
- Create a new workspace,
- Now, pay attention to the License Type

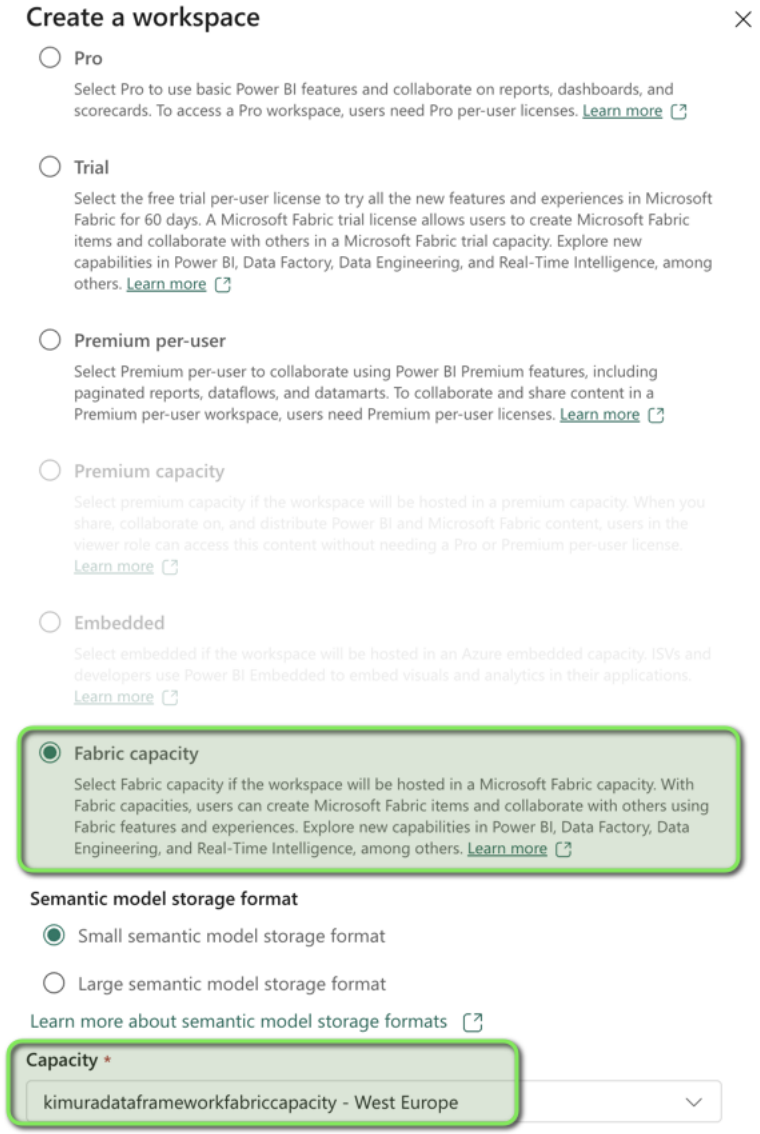

Creating the first Fabric Workspace

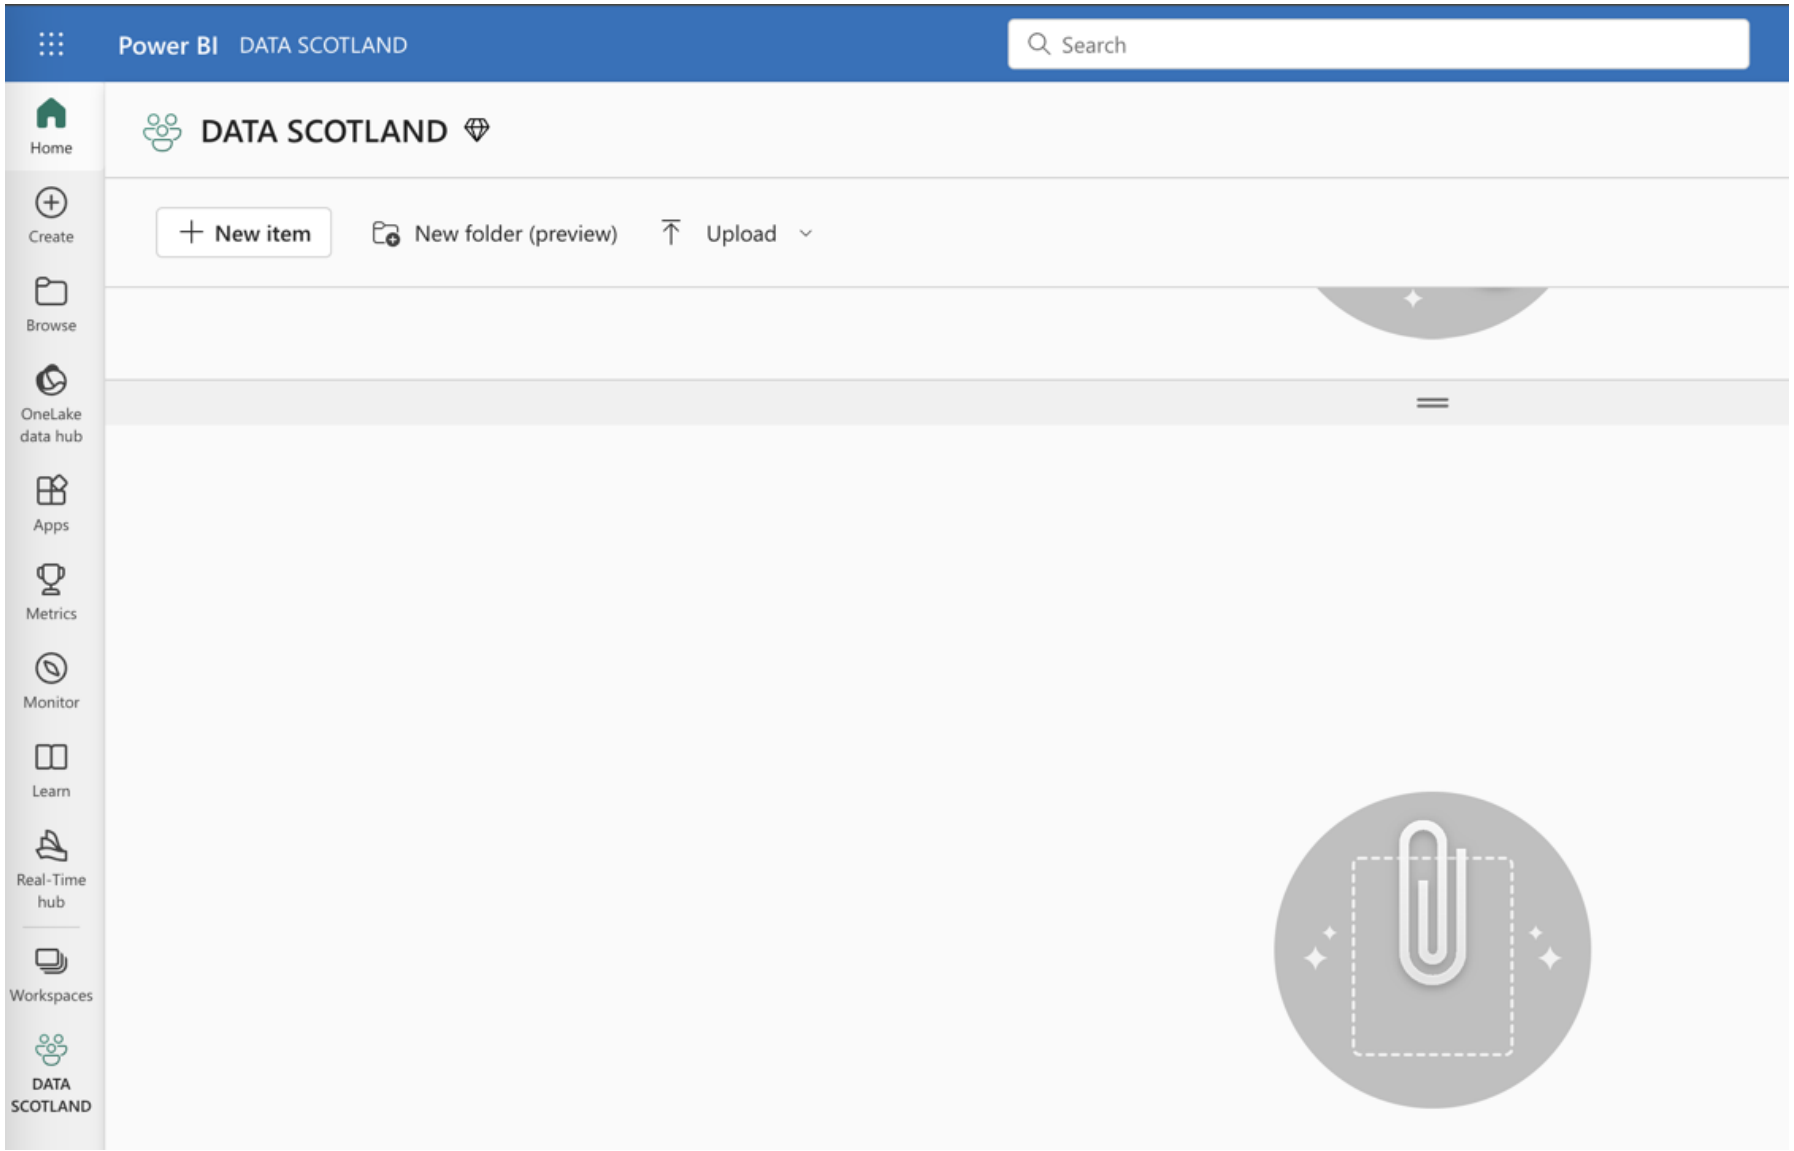

Medallion architecture

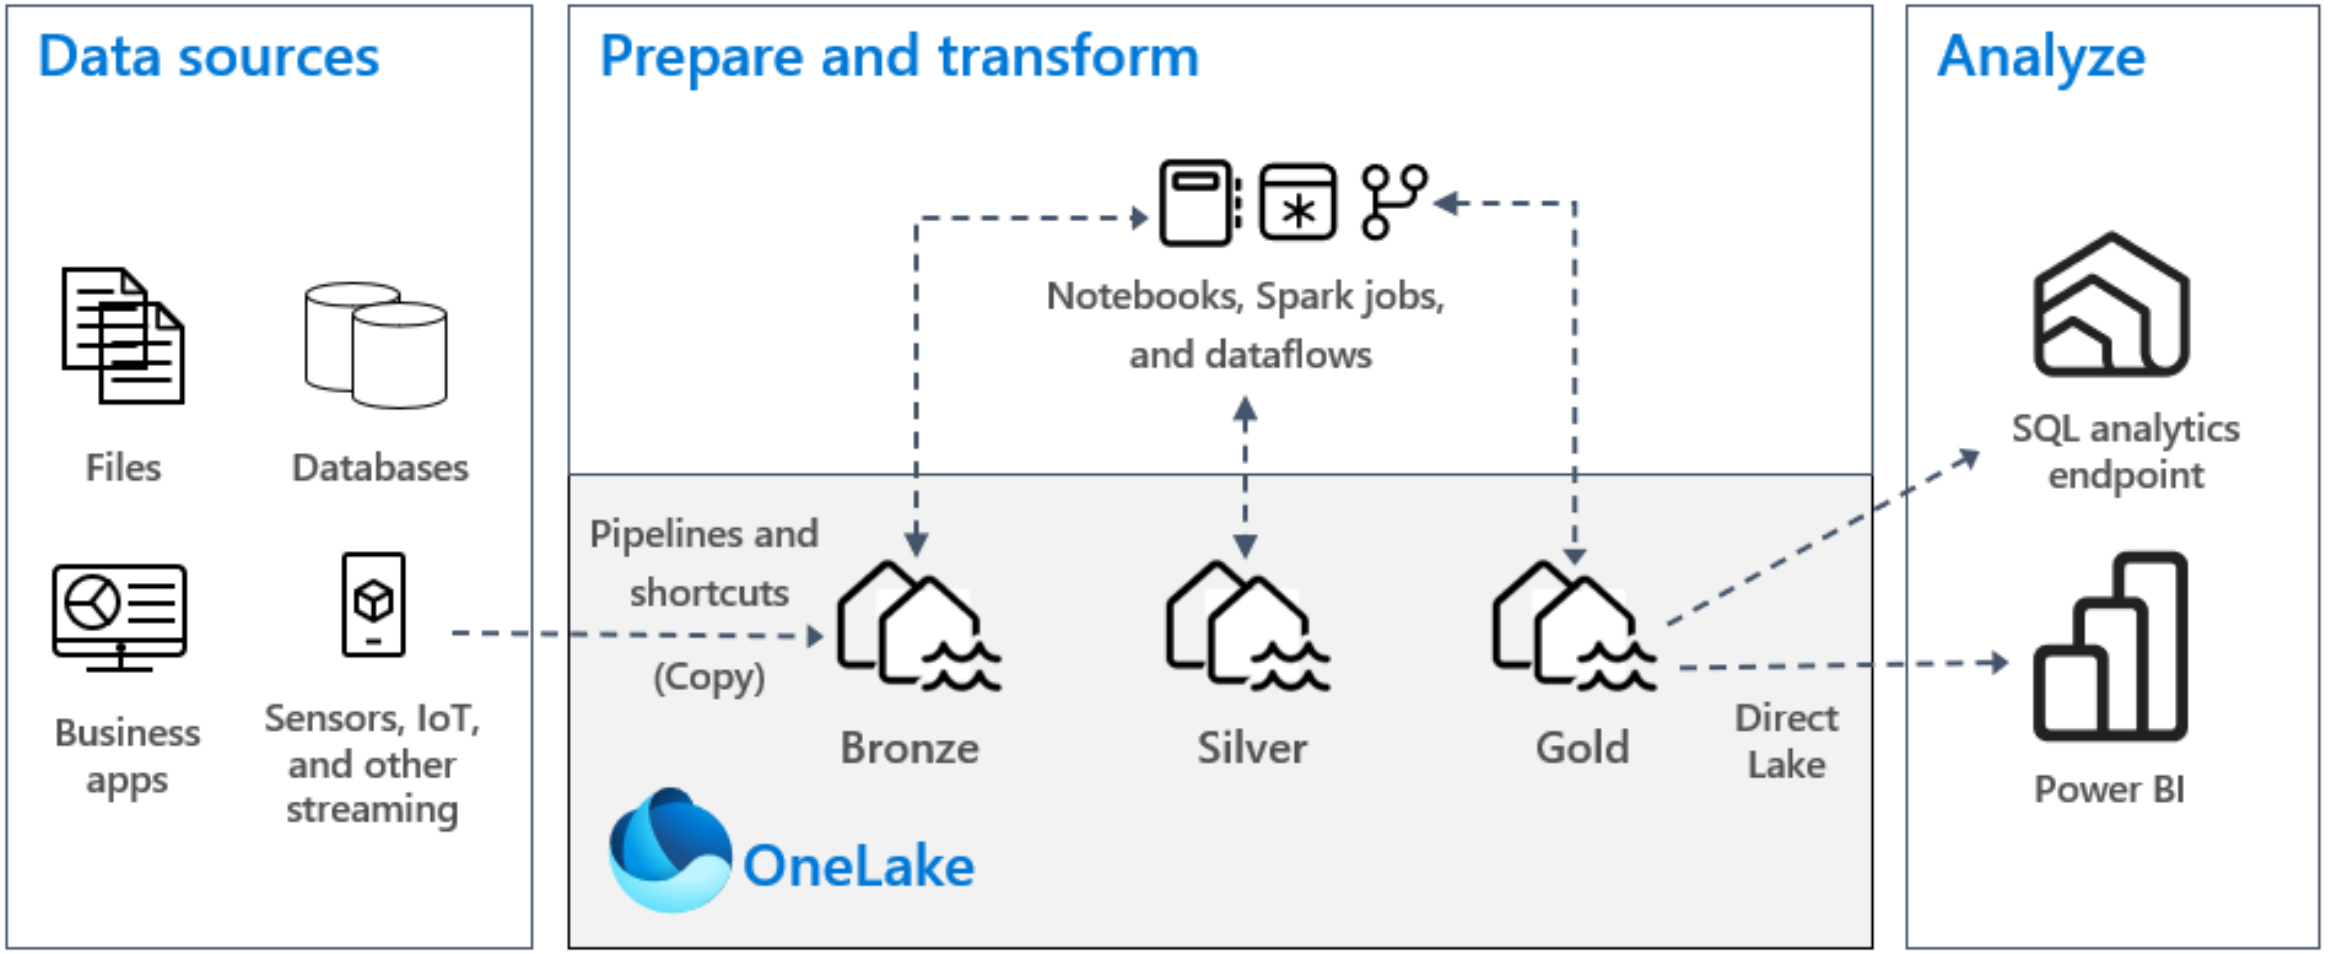

Workspaces and Medallion

- One workspace per Medallion stage

- One workspace per DTAP environment

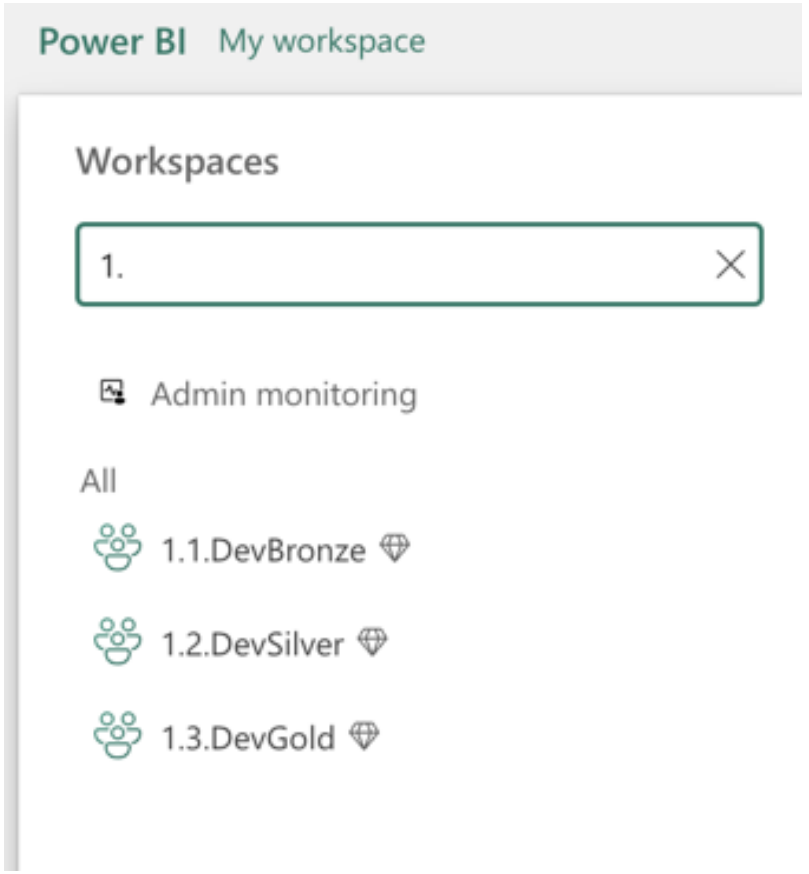

Your first Fabric Lakehouse

The scenario

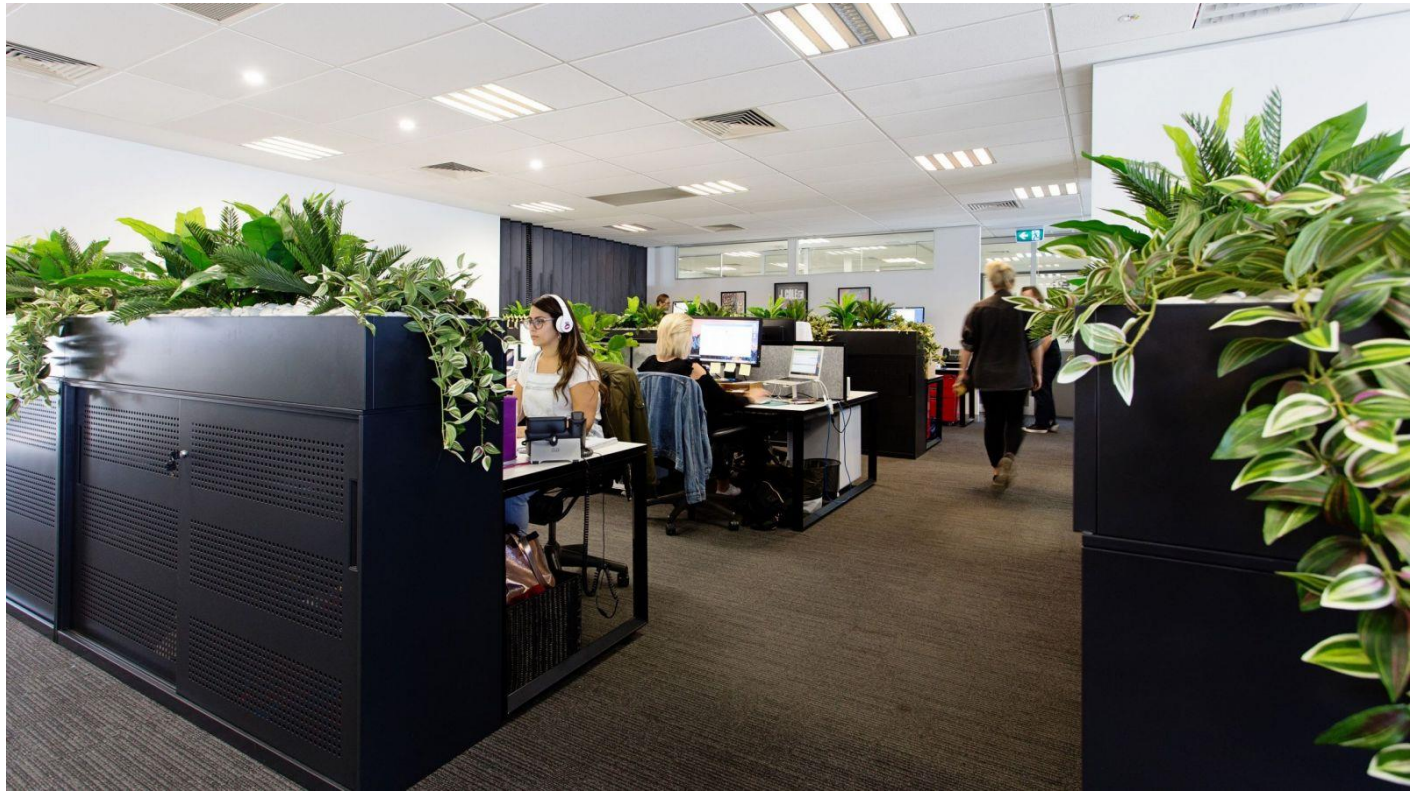

- Client: small B2B professional services company
- Main software: Moneybird, Odoo
- Data domains: Finance, Projects & Timesheets

The rulebook

- Data Factory Pipelines orchestration
- PySpark notebooks for processing
- Reusable code for everything
- Lakehouse for storage

Prerequisites

Setting up:

- Fabric compute capacity (done)
- Fabric workspaces (done)
- Configuring Spark (we skip that for today)
- Setting up the Environment (**important**!)

Environments in Fabric

- Fabric contains a lot of Python packages natively
- You can also add packages using the PyPI (Python Package Index)
- And, you can write and include your own code
- This is done using Environments

Environments in Fabric

- Create a new environment:

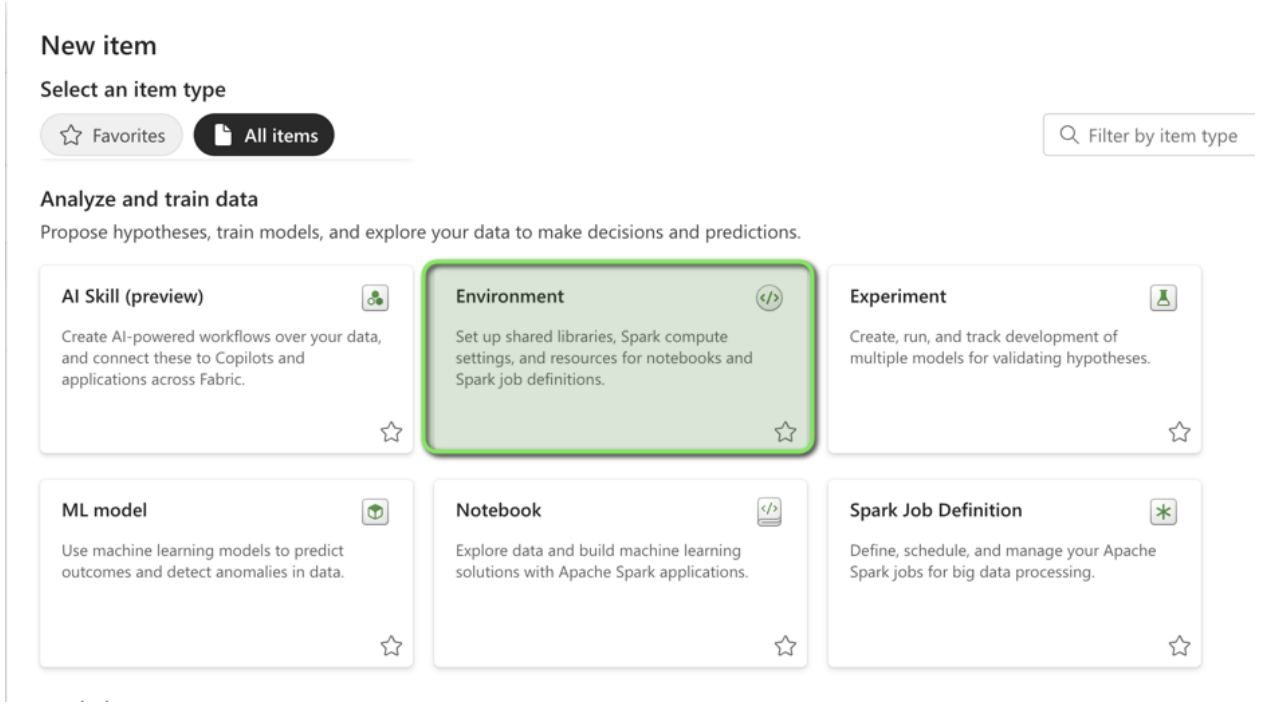

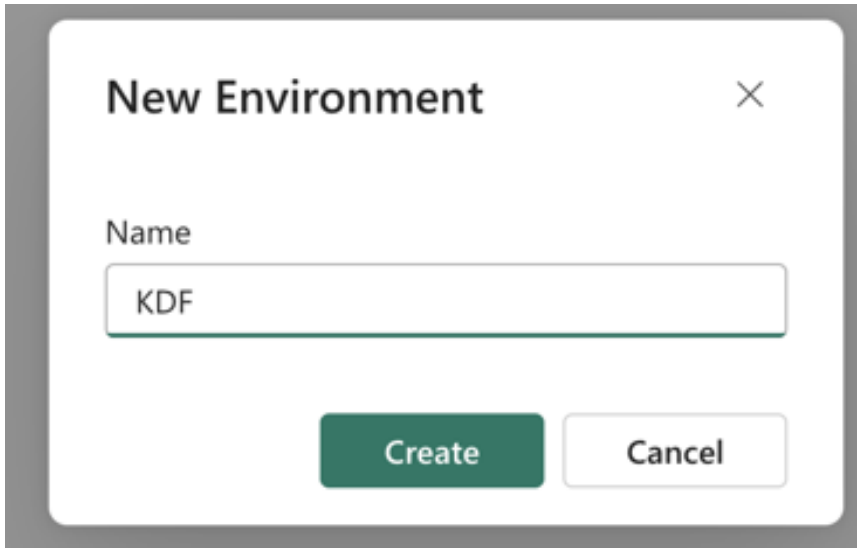

Environments in Fabric

- Add code from PyPI:

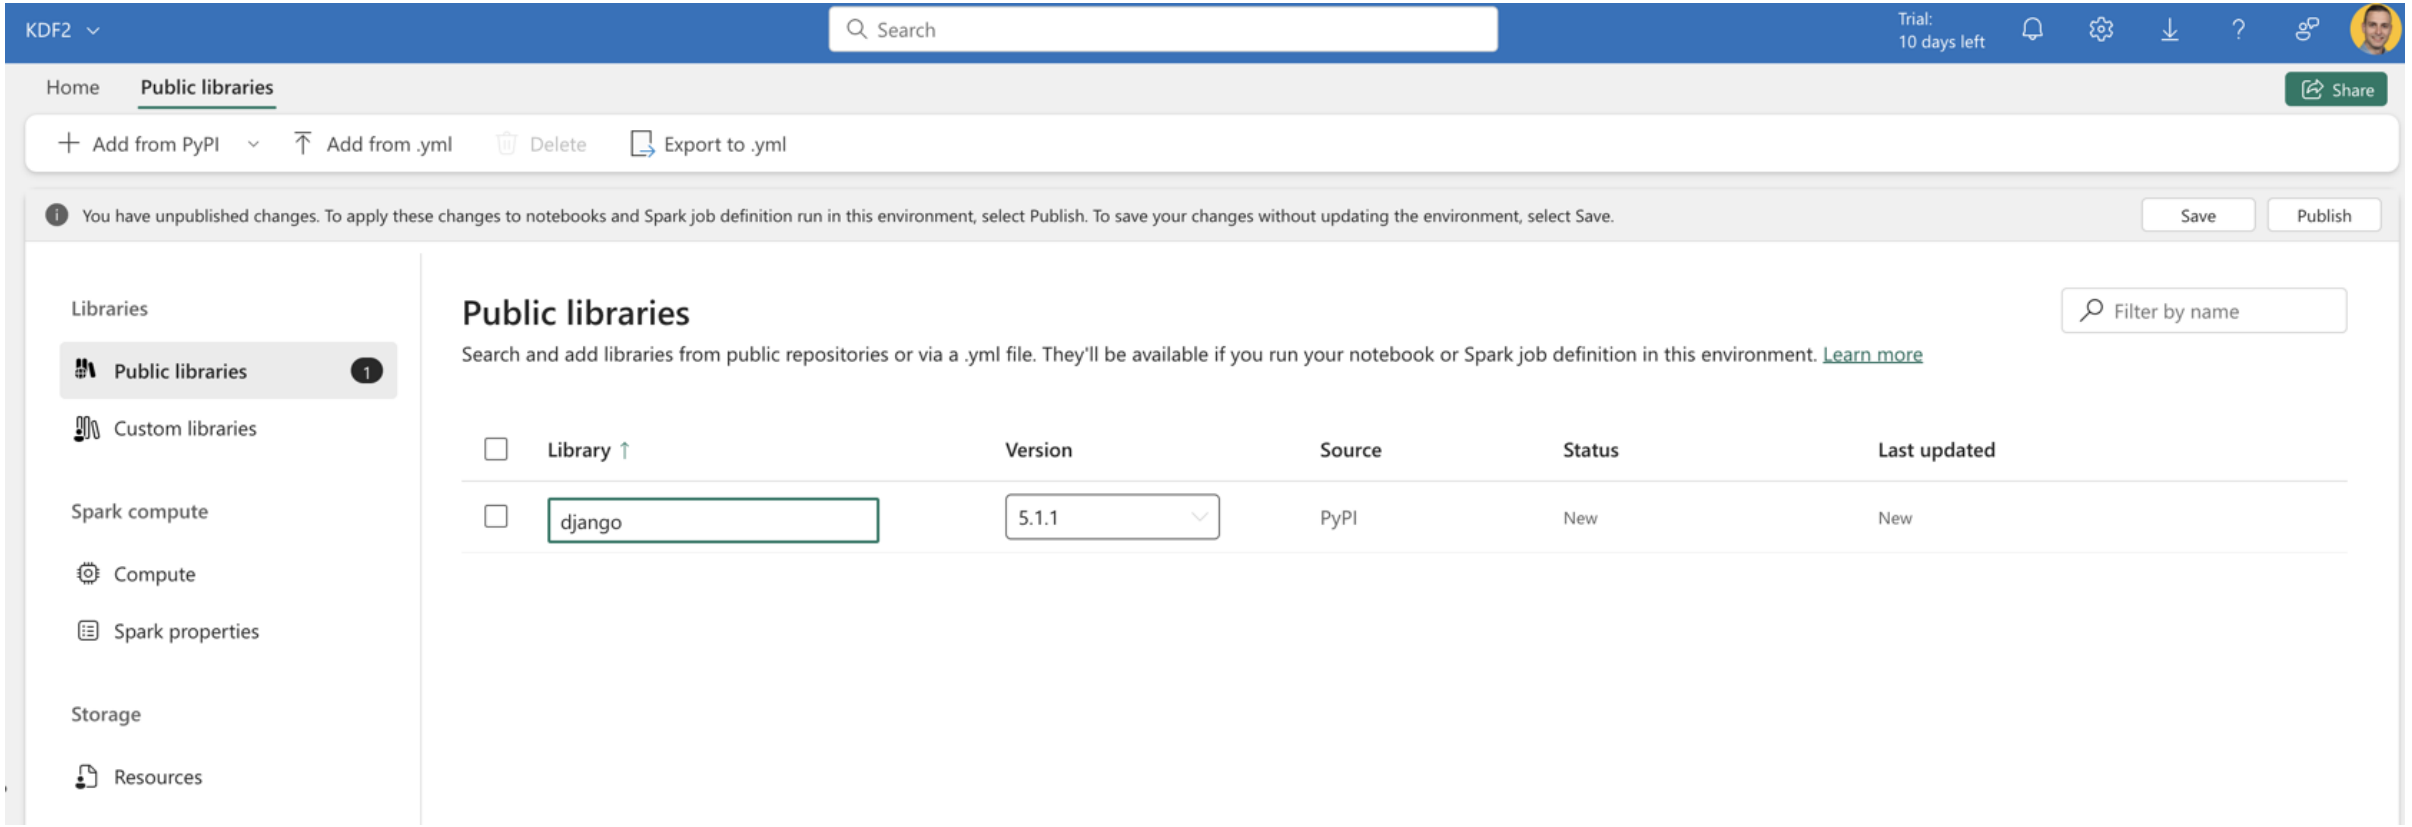

Environments in Fabric

- Add your own code:

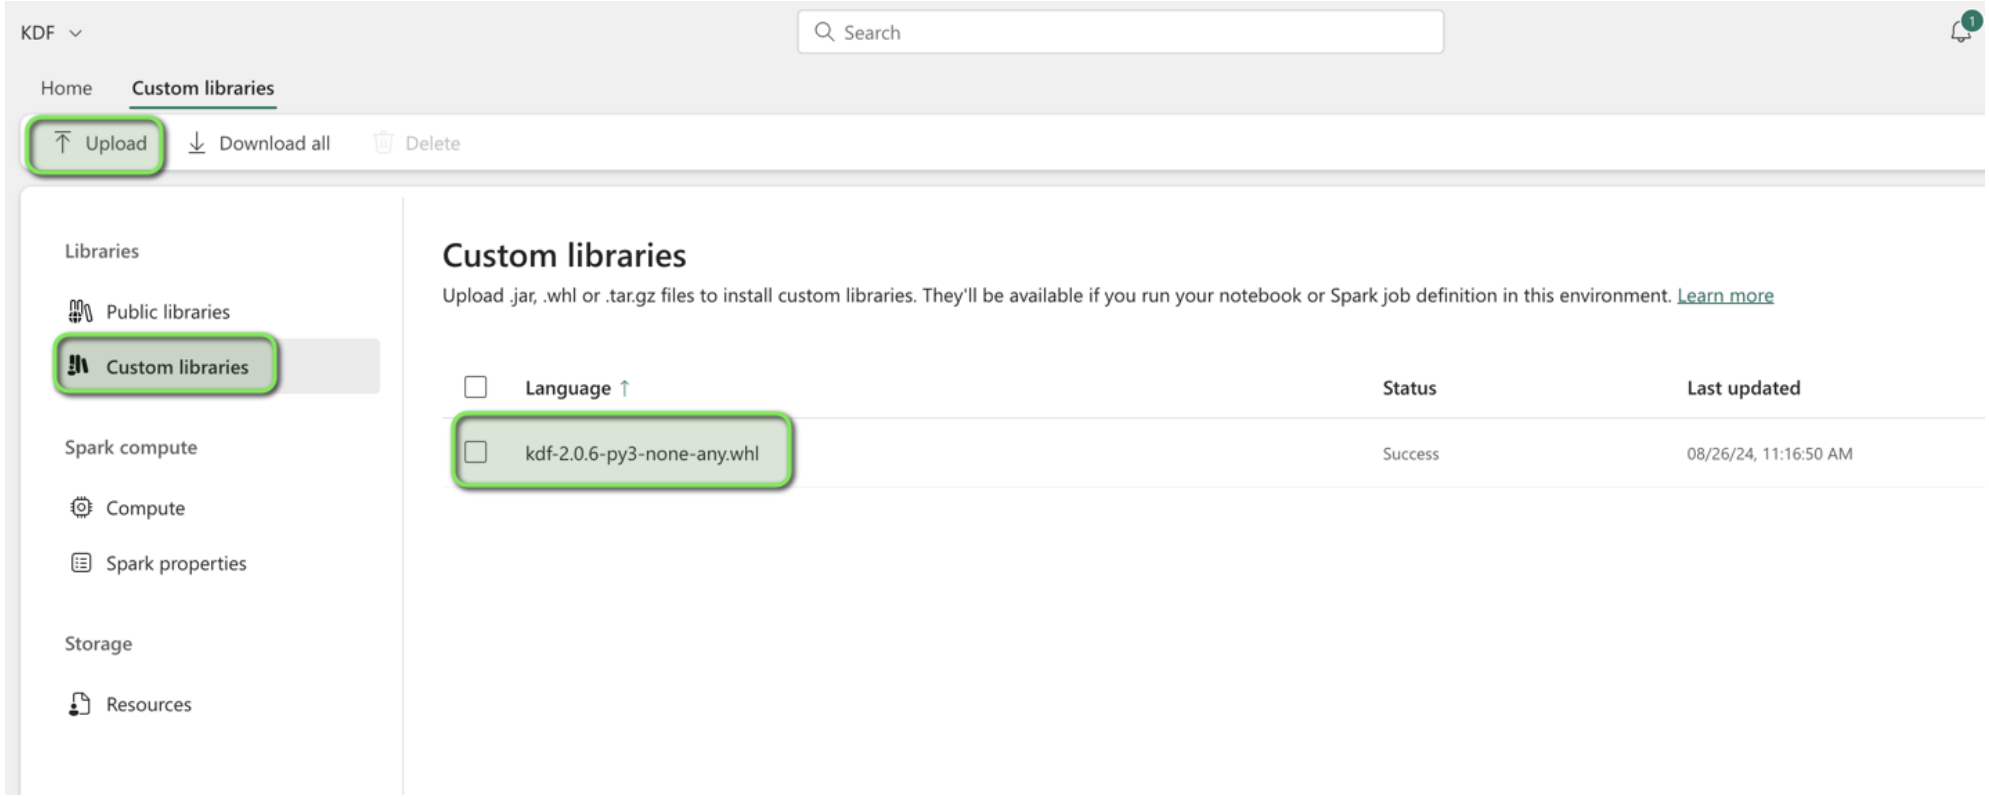

Three lakehouses

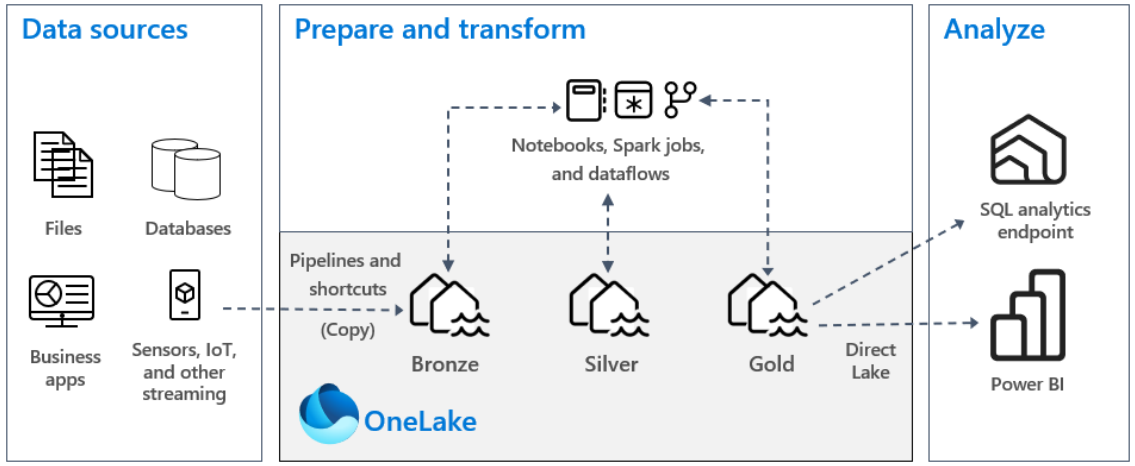

- Bronze: timestamped folders, (incremental loads), raw data (files; json, xml, csv, parquet etc)
- Silver: historical archive, normalised schemas, delta tables
- Gold: prepared and modeled DWH tables, dims and facts, delta tables

Setting up the ELT process

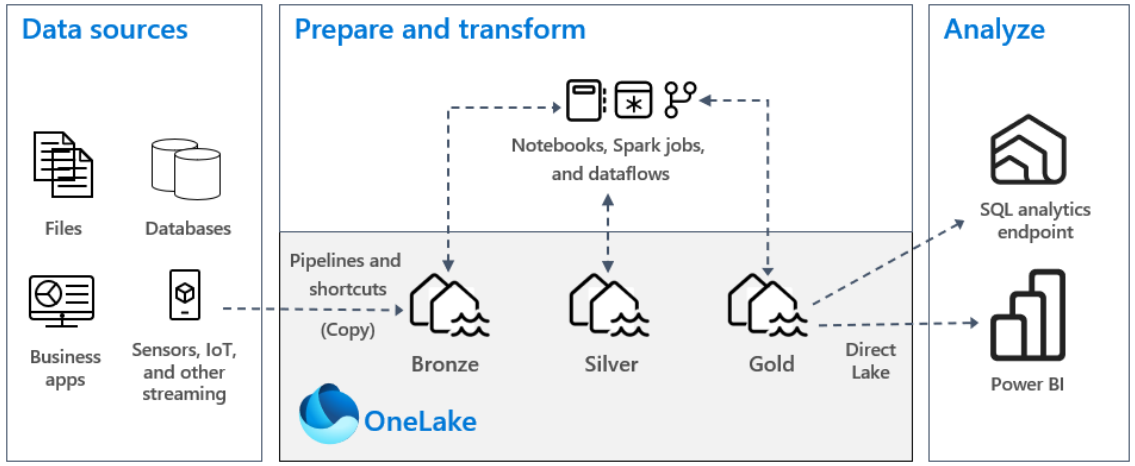

- Keep in mind the Medallion architecture
- Start connecting to data, to extract Moneybird and Odoo to Bronze
- We do so using Python, executed in Fabric Notebooks, scheduled and run in Fabric Pipelines

ELT: Extract

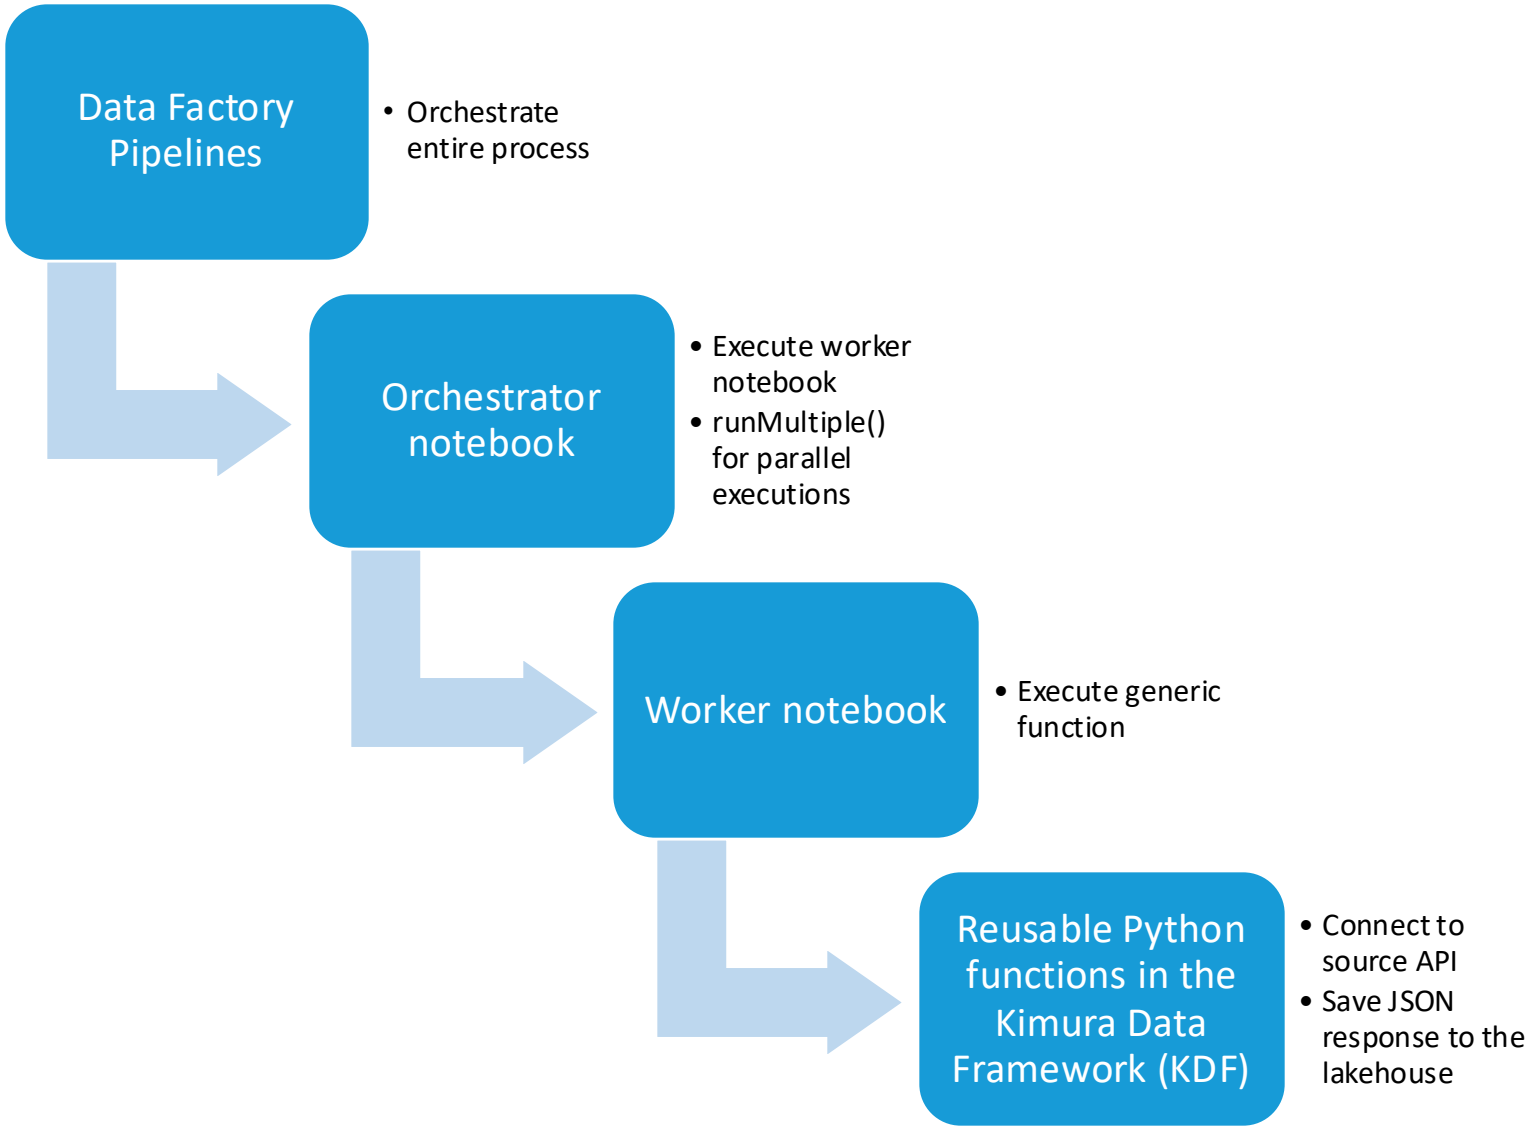
Data Factory Pipelines
• Orchestrate entire process
Orchestrator notebook
• Execute worker notebook
• runMultiple() for parallel executions
Worker notebook
• Execute generic function
Reusable Python functions in the Kimura Data Framework (KDF)
• Connect to source API
• Save JSON response to the lakehouse

ELT: Bronze folder structure

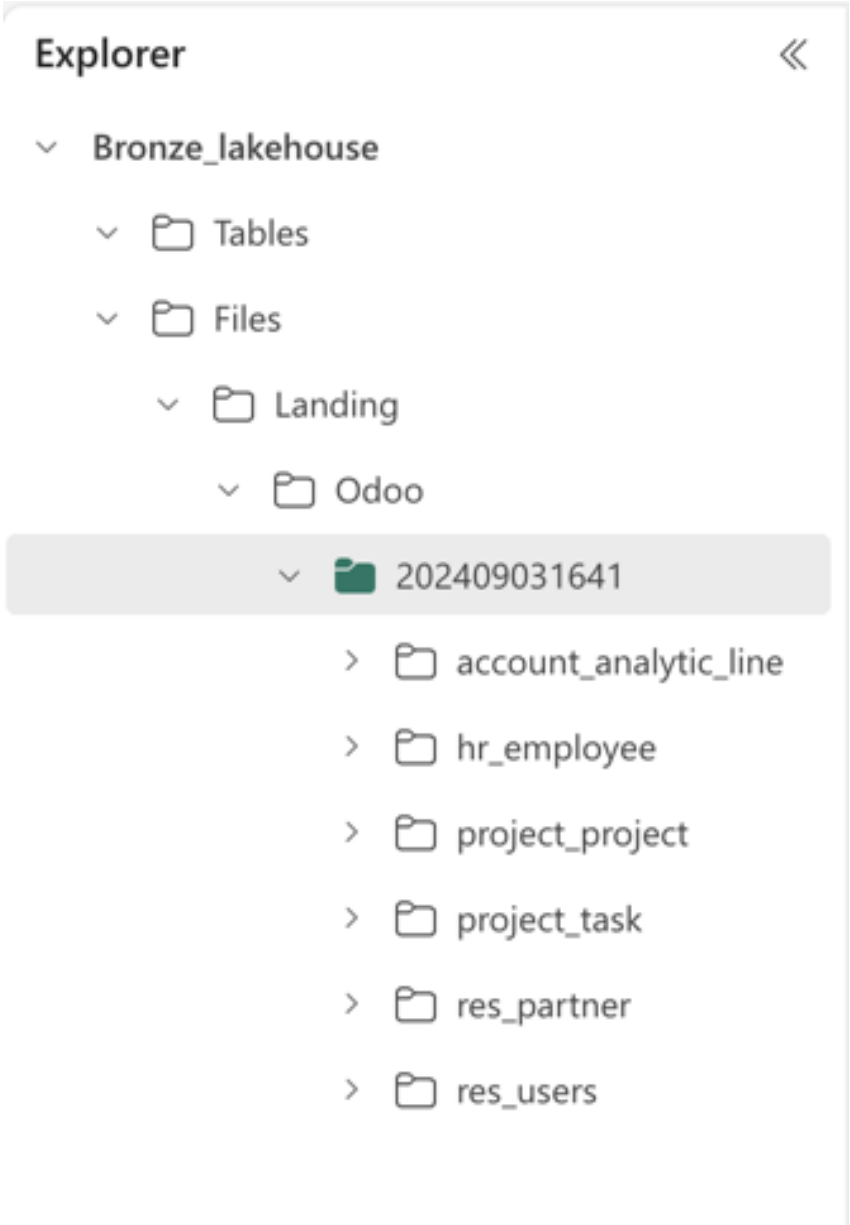

- Root folder ‘Landing’ (for shortcuts!)
- Source system
- Timestamp
- Folder per table/endpoint
- Files managed by Fabric!

Silver lakehouse

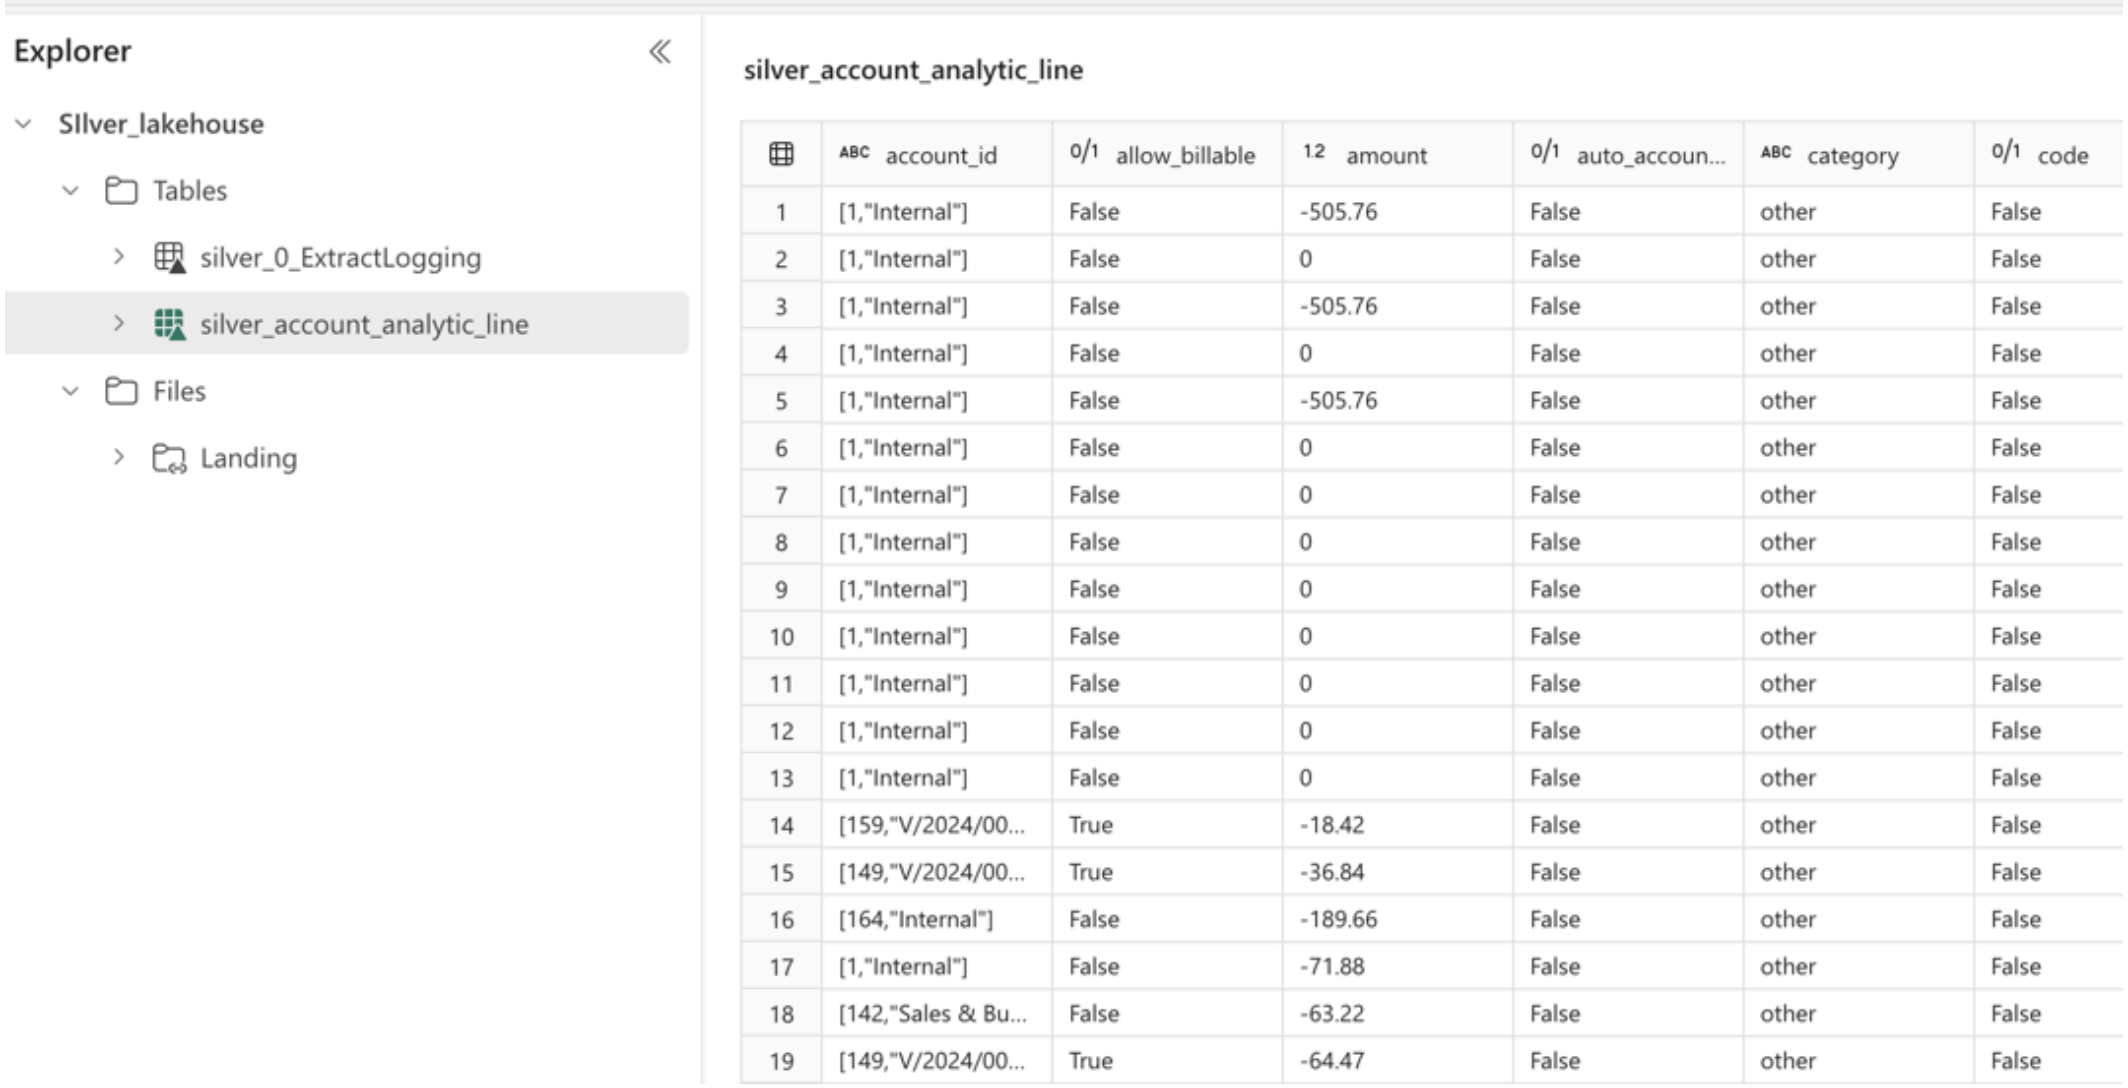

silver_account_analytic_line

	account_id	allow_billable	amount	auto_accoun...	category	code
1	[1,"Internal"]	False	-505.76	False	other	False
2	[1,"Internal"]	False	0	False	other	False
3	[1,"Internal"]	False	-505.76	False	other	False
4	[1,"Internal"]	False	0	False	other	False
5	[1,"Internal"]	False	-505.76	False	other	False
6	[1,"Internal"]	False	0	False	other	False
7	[1,"Internal"]	False	0	False	other	False
8	[1,"Internal"]	False	0	False	other	False
9	[1,"Internal"]	False	0	False	other	False
10	[1,"Internal"]	False	0	False	other	False
11	[1,"Internal"]	False	0	False	other	False
12	[1,"Internal"]	False	0	False	other	False
13	[1,"Internal"]	False	0	False	other	False
14	[159,"V/2024/00...	True	-18.42	False	other	False
15	[149,"V/2024/00...	True	-36.84	False	other	False
16	[164,"Internal"]	False	-189.66	False	other	False
17	[1,"Internal"]	False	-71.88	False	other	False
18	[142,"Sales & Bu...	False	-63.22	False	other	False
19	[149,"V/2024/00...	True	-64.47	False	other	False

Gold lakehouse

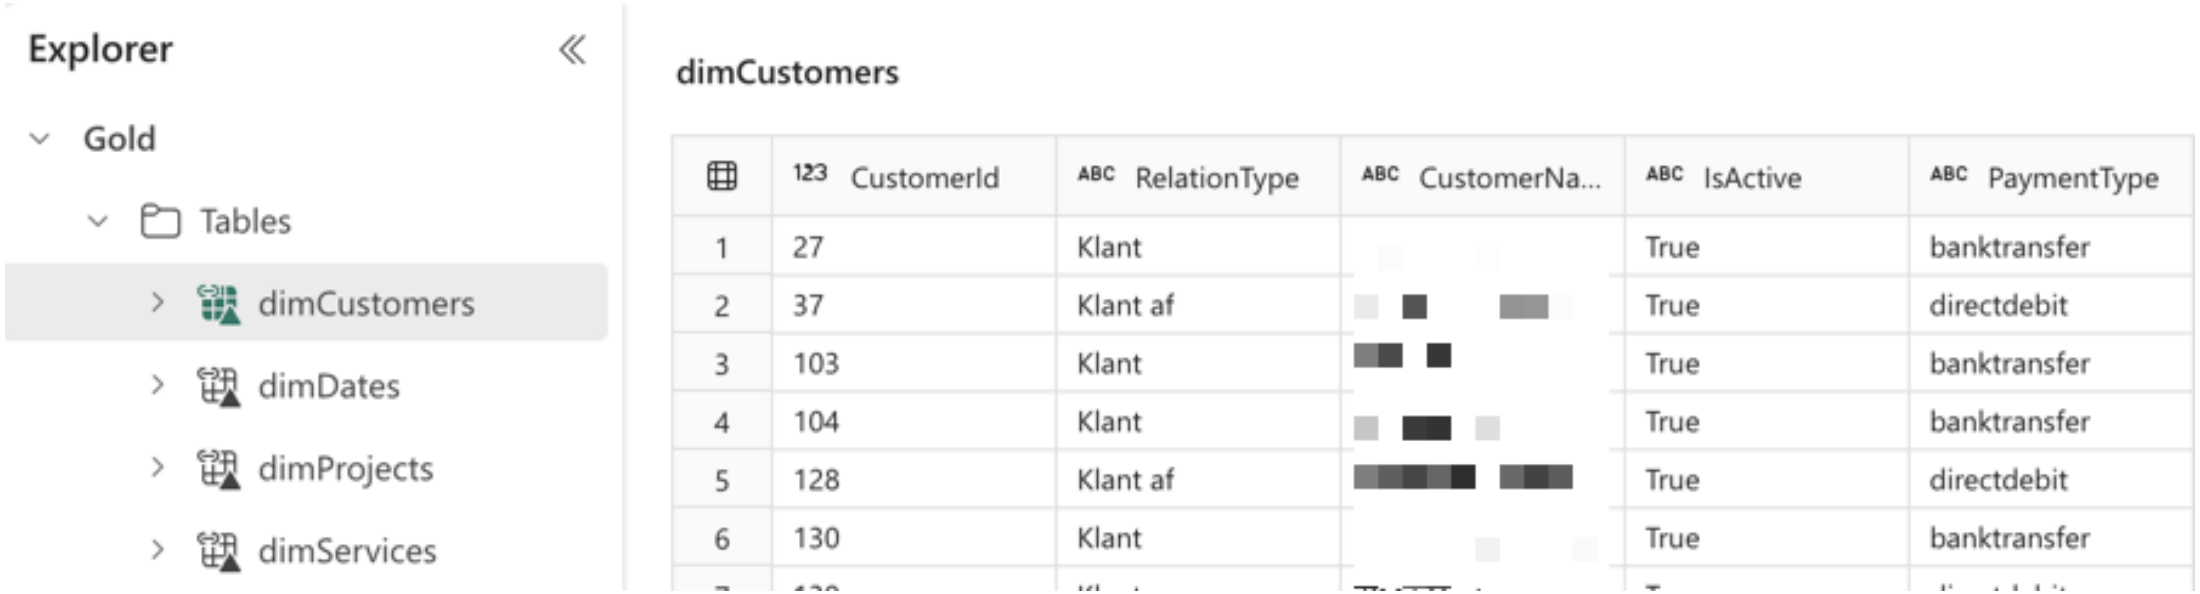

End-to-end solution with examples

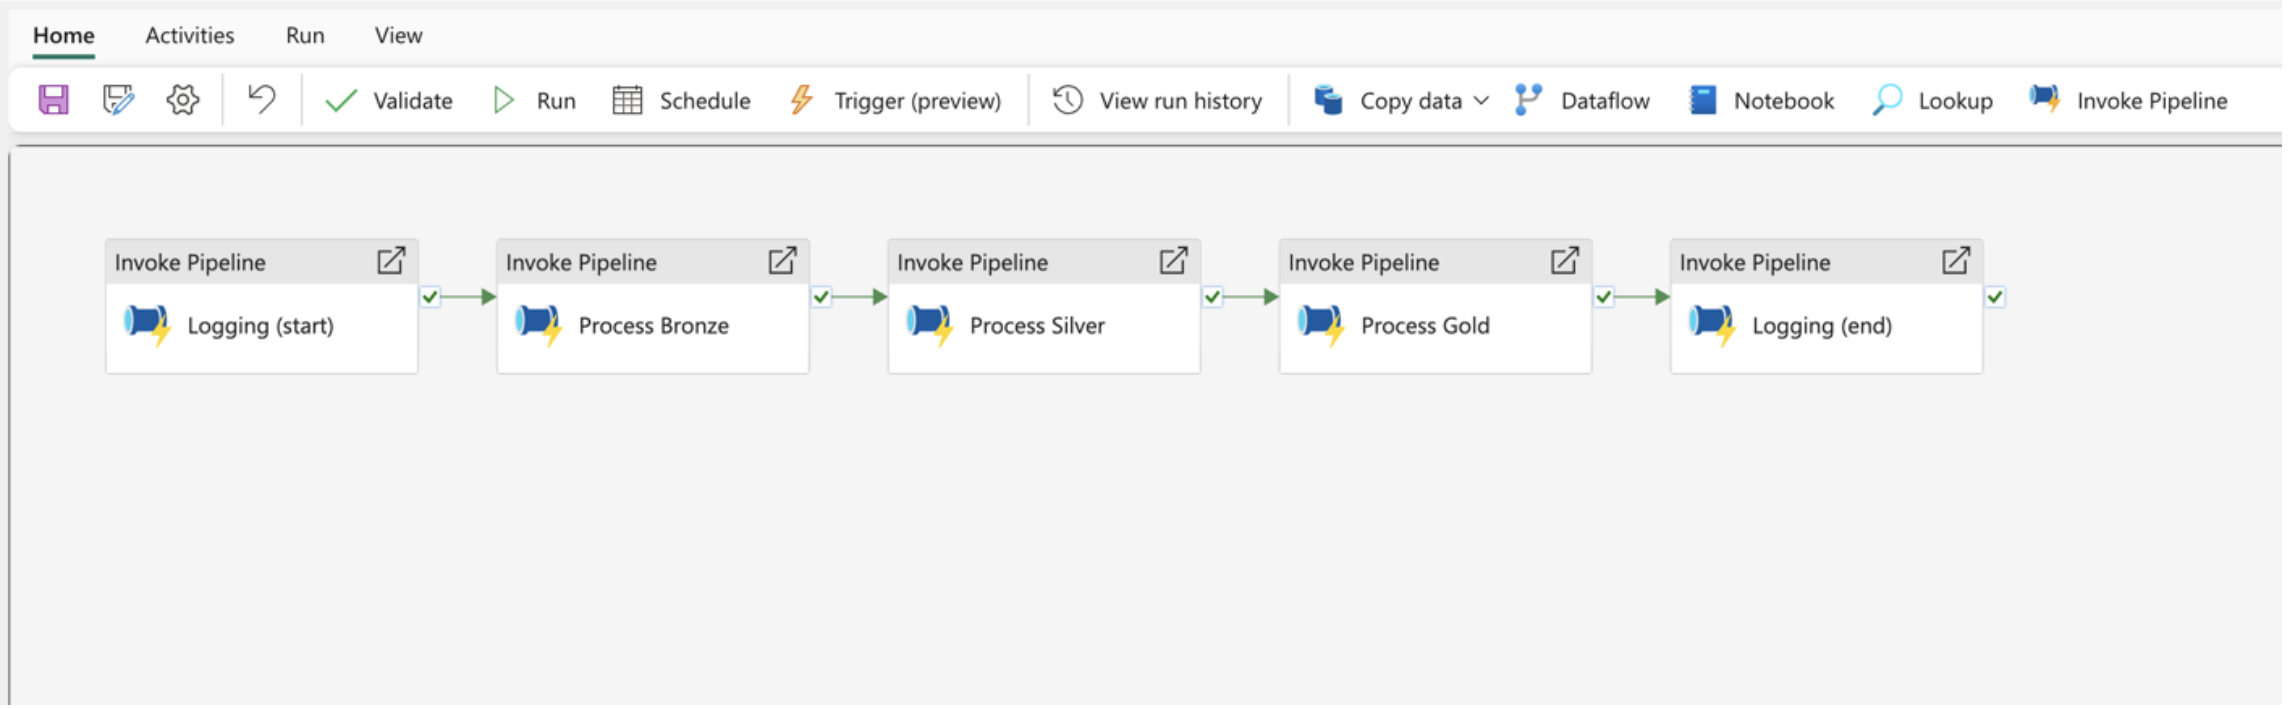

End-to-end solution with examples

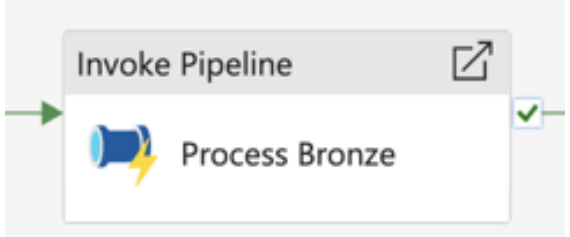

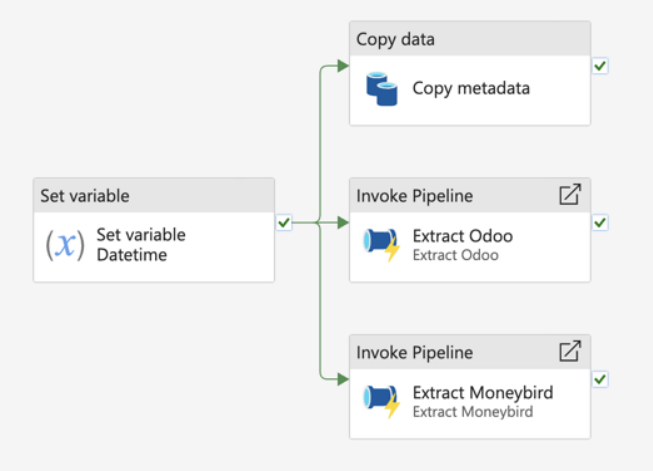

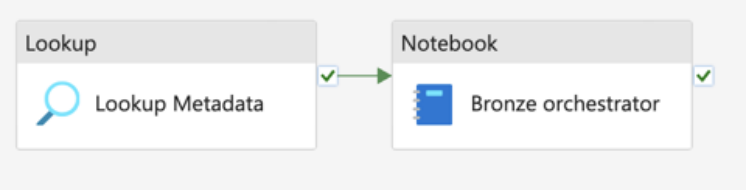

End-to-end solution with examples

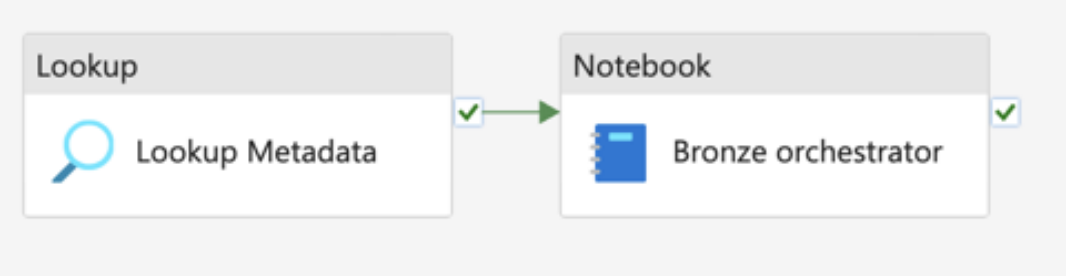


```
for index,
    endpoi

    for adr
        en

        la

        ex

        if

        if

        en

    api_en

    activi
        'na
        'pa
        't
        'a

        }
    }

    activi

DAG = {
    "activ
}

mssparkuti

new_timesta
new_timesta
```

```
    api_endpoints = ','.join(endpoints)

    activity = {
        'name': row['output_filename'],
        'path': 'MB extract GET',
        'timeoutPerCellInSeconds': 9000,
        'args': {
            'output_filename': row['output_filename'],
            'timestamp': timestamp,
            'api_endpoints': api_endpoints,
            'api_tokens': api_tokens
        }
    }

    activities.append(activity)

DAG = {
    "activities": activities
}

mssparkutils.notebook.runMultiple(DAG)

new_timestamps = [Row(administration_id=int(admin_id), tir
```

```
1   if output_filename == 'financial_mutations':
2       mb.get_financial_mutations(
3           api_endpoints=api_endpoints.split(','),
4           api_tokens=api_tokens.split(','),
5           landing_path=landingzone_path,
6           output_filename=output_filename
7       )
8   else:
9       mb.get_endpoint(
10          api_endpoints=api_endpoints.split(','),
11          api_tokens=api_tokens.split(','),
12          landing_path=landingzone_path,
13          output_filename=output_filename
14      )
```

- Command executed in 24 sec 371 ms by Kevin Land on 11:12:50 AM, 8/26/24

End-to-end solution with examples

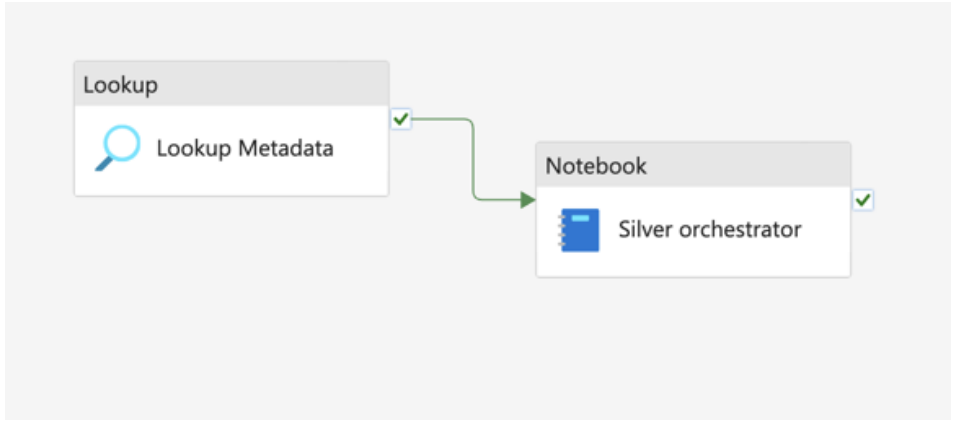

```python
#Mapping logic
gold_table = "dimServices"
business_key_columns = ['ServiceId']

df_mapping = spark.sql("""
select
    t.`data.id` as ServiceId
    ,t.`data.project_id` as ProjectId
    ,t.`data.name` as ServiceName
    ,t.`data.invoice_method` as InvoiceMethod
    ,t.`data.amount` as Amount
    ,t.`data.start_date` as StartDate
    ,t.`data.end_date` as EndDate
    ,t.`data.revenue_group.label` as RevenueGroup
    ,t.`data.subscription_cycle` as InvoiceCycle
    ,t.`data.price` as Price
    ,t.`data.amount` * t.`data.price` as Value
    ,case when t.`data.revenue_group.id` = 'revenuegroup:ba06e5bec6c1b6b5' then 'Licenses' else 'Other' end as MRRLabel
    ,c.Sys_ID as FK_dimCustomers_Sys_ID
from Bronze.bronze_Odoo_ProjectsService t
left join Bronze.bronze_Odoo_ProjectsProject p
    on t.`data.project_id` = p.`data.id`
left join Silver.silver_dimCustomers c
    on p.`data.organization.id` = c.OdooId
where 1=1
""")

k.gold_load_table( \
    gold_table = gold_table_name, \
    df_mapping = df_mapping, \
    business_key_columns = business_key_columns, \
    spark_session = spark\
)
```

Recap

Recap

1. Keep core-concepts in mind

2. Setup a rulebook and abide by it

3. Use tools and reusable code!

4. Fabric lakehouses aren’t that scary ☺

Bas Land

- BI consultant since 2013
- Co-founder of two data companies:
  - Kimura Data Intelligence (consultancy)
  - DataChimp (SaaS analytics for accounting firms)

- Married to Anouk, we have a daxhund (☺) called Chester
- Sports: purple belt Brazilian jiu-jitsu, weight lifting, running

Kimura Data Intelligence B.V.

Fonteinkruid 6b
3931 WX, Woudenberg
The Netherlands

www.kimura.nl
info@kimura.nl

Bas Land
Managing Partner

bas@kimura.nl

Kimura
data intelligence

DATA:Scotland 2024

Session Feedback

Event Feedback