

Fabric Lakehouse

for SQL Data Engineers



Bas Land



- BI consultant since 2013
- Co-founder of two data companies:
 - Kimura Data Intelligence (consultancy)
 - DataChimp (SaaS analytics for accounting firms)
- Married to Anouk, we have a dachshund (😊) called Chester
- Sports: Brazilian jiu-jitsu, weight lifting, running



How do you manage all of that?



Data warehouse experience



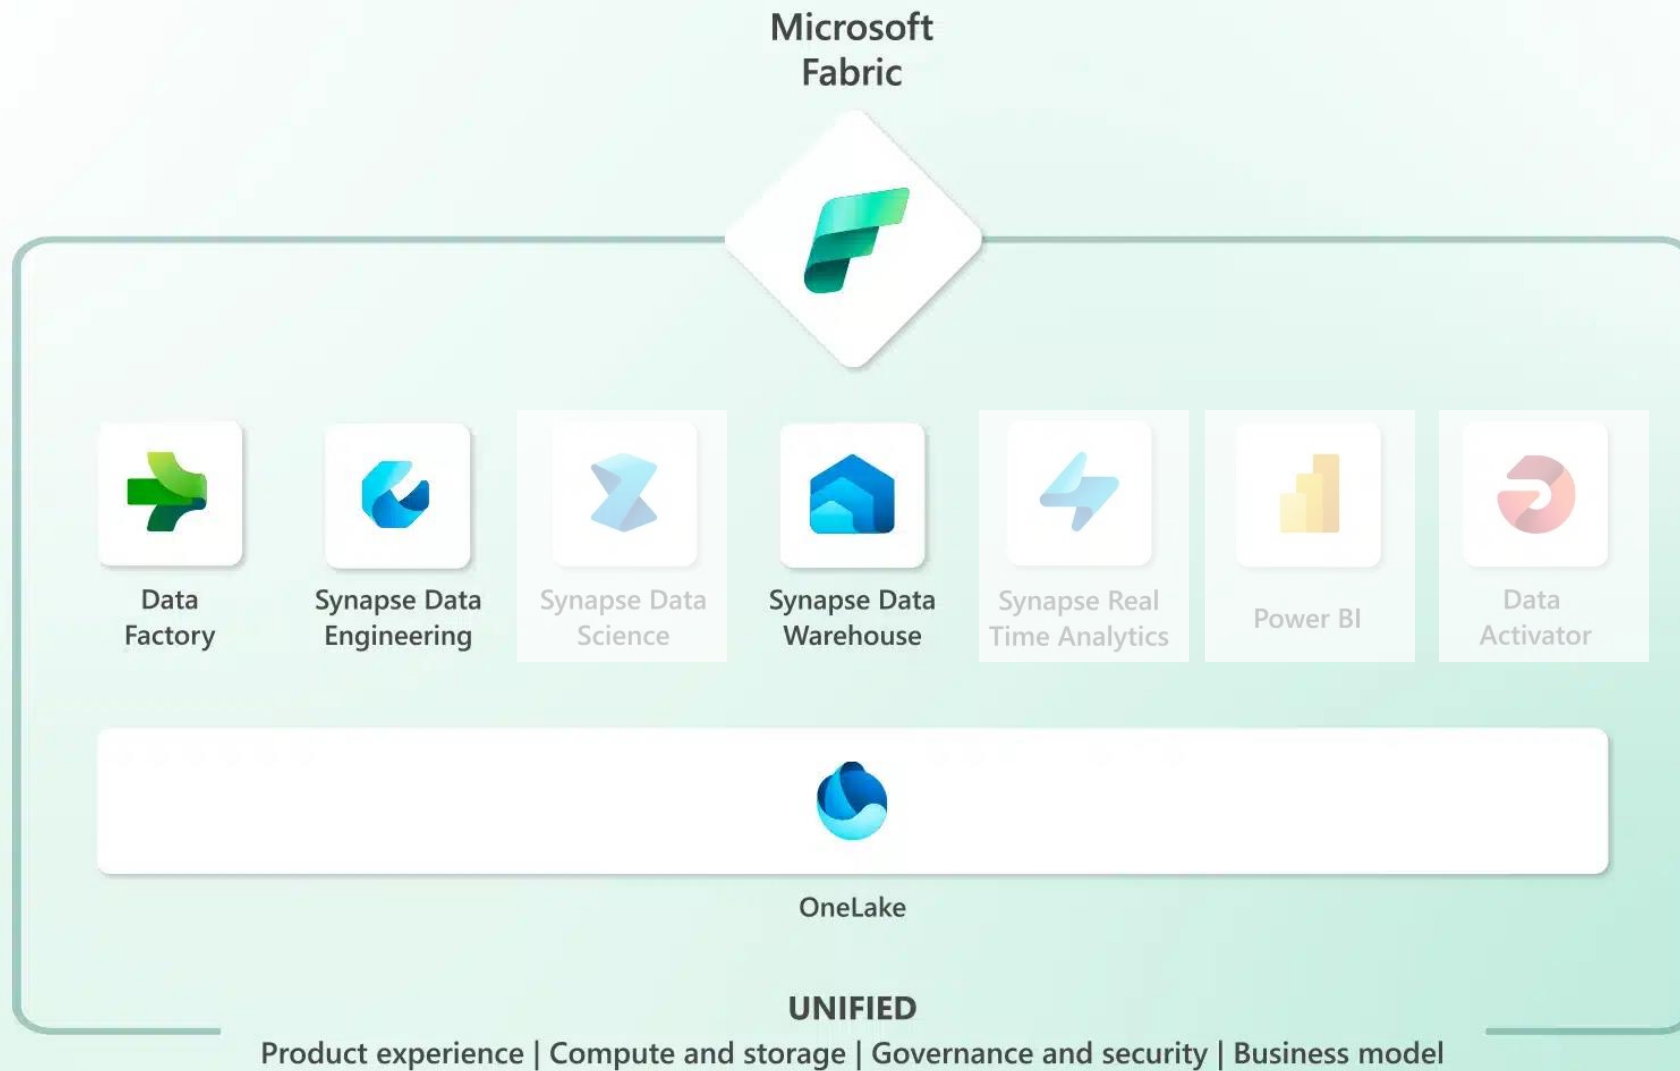
- Since 2013 I have personally built >20 data warehouses, guided >50 in total
- Most in SQL technology (SQL Server and Azure SQL)
- And now we've done 4 MS Fabric implementations with 2 more on the way
- Moving from SQL to Fabric Lakehouses is easier than it seems!

Today



- MS Fabric for SQL Data Engineers!
 1. Data warehouse design patterns
 2. Battle of the engines: SQL vs Spark
 3. Code reusability
- Your first data warehouse with Fabric

What the F?



Data warehouse – general design principles

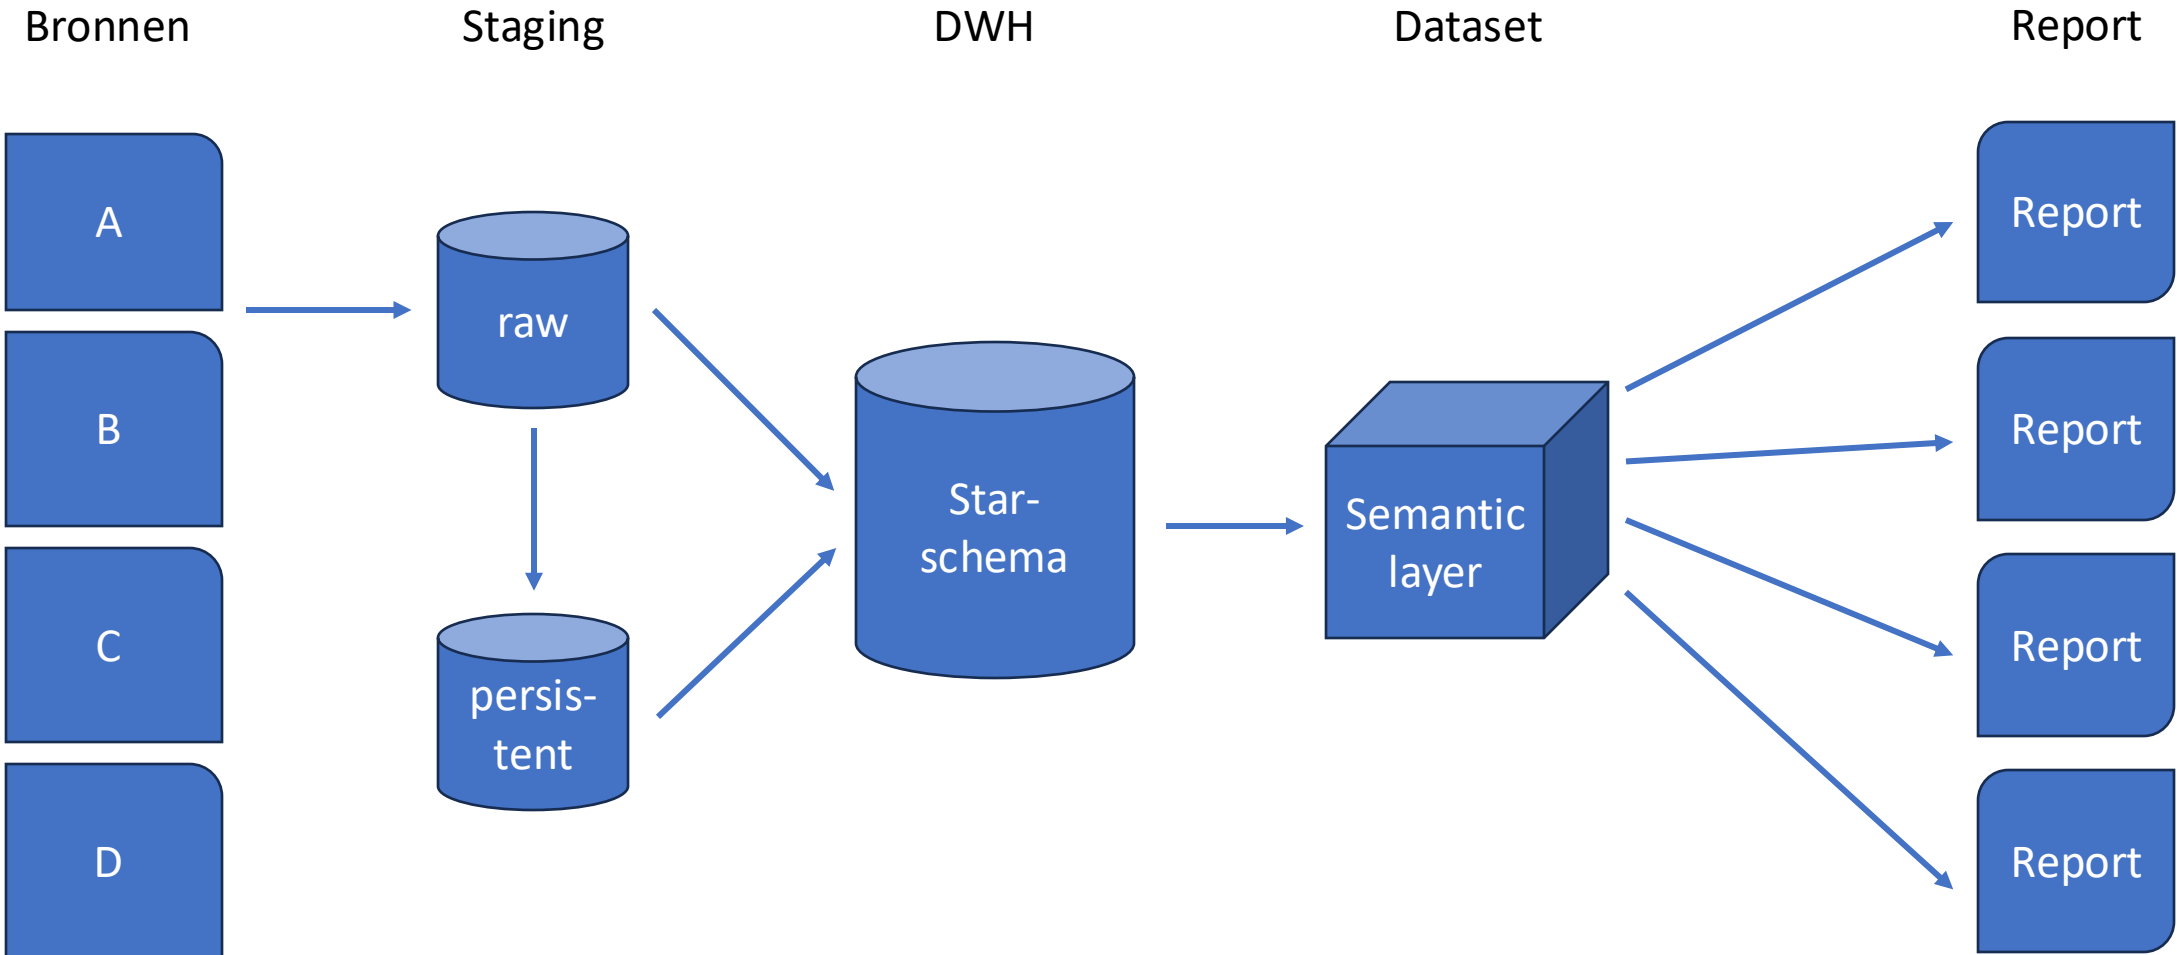
- Historical data archive from operational applications
- Combined data from multiple applications
- Prepared / precalculated transformations and business logic
- Single source of truth



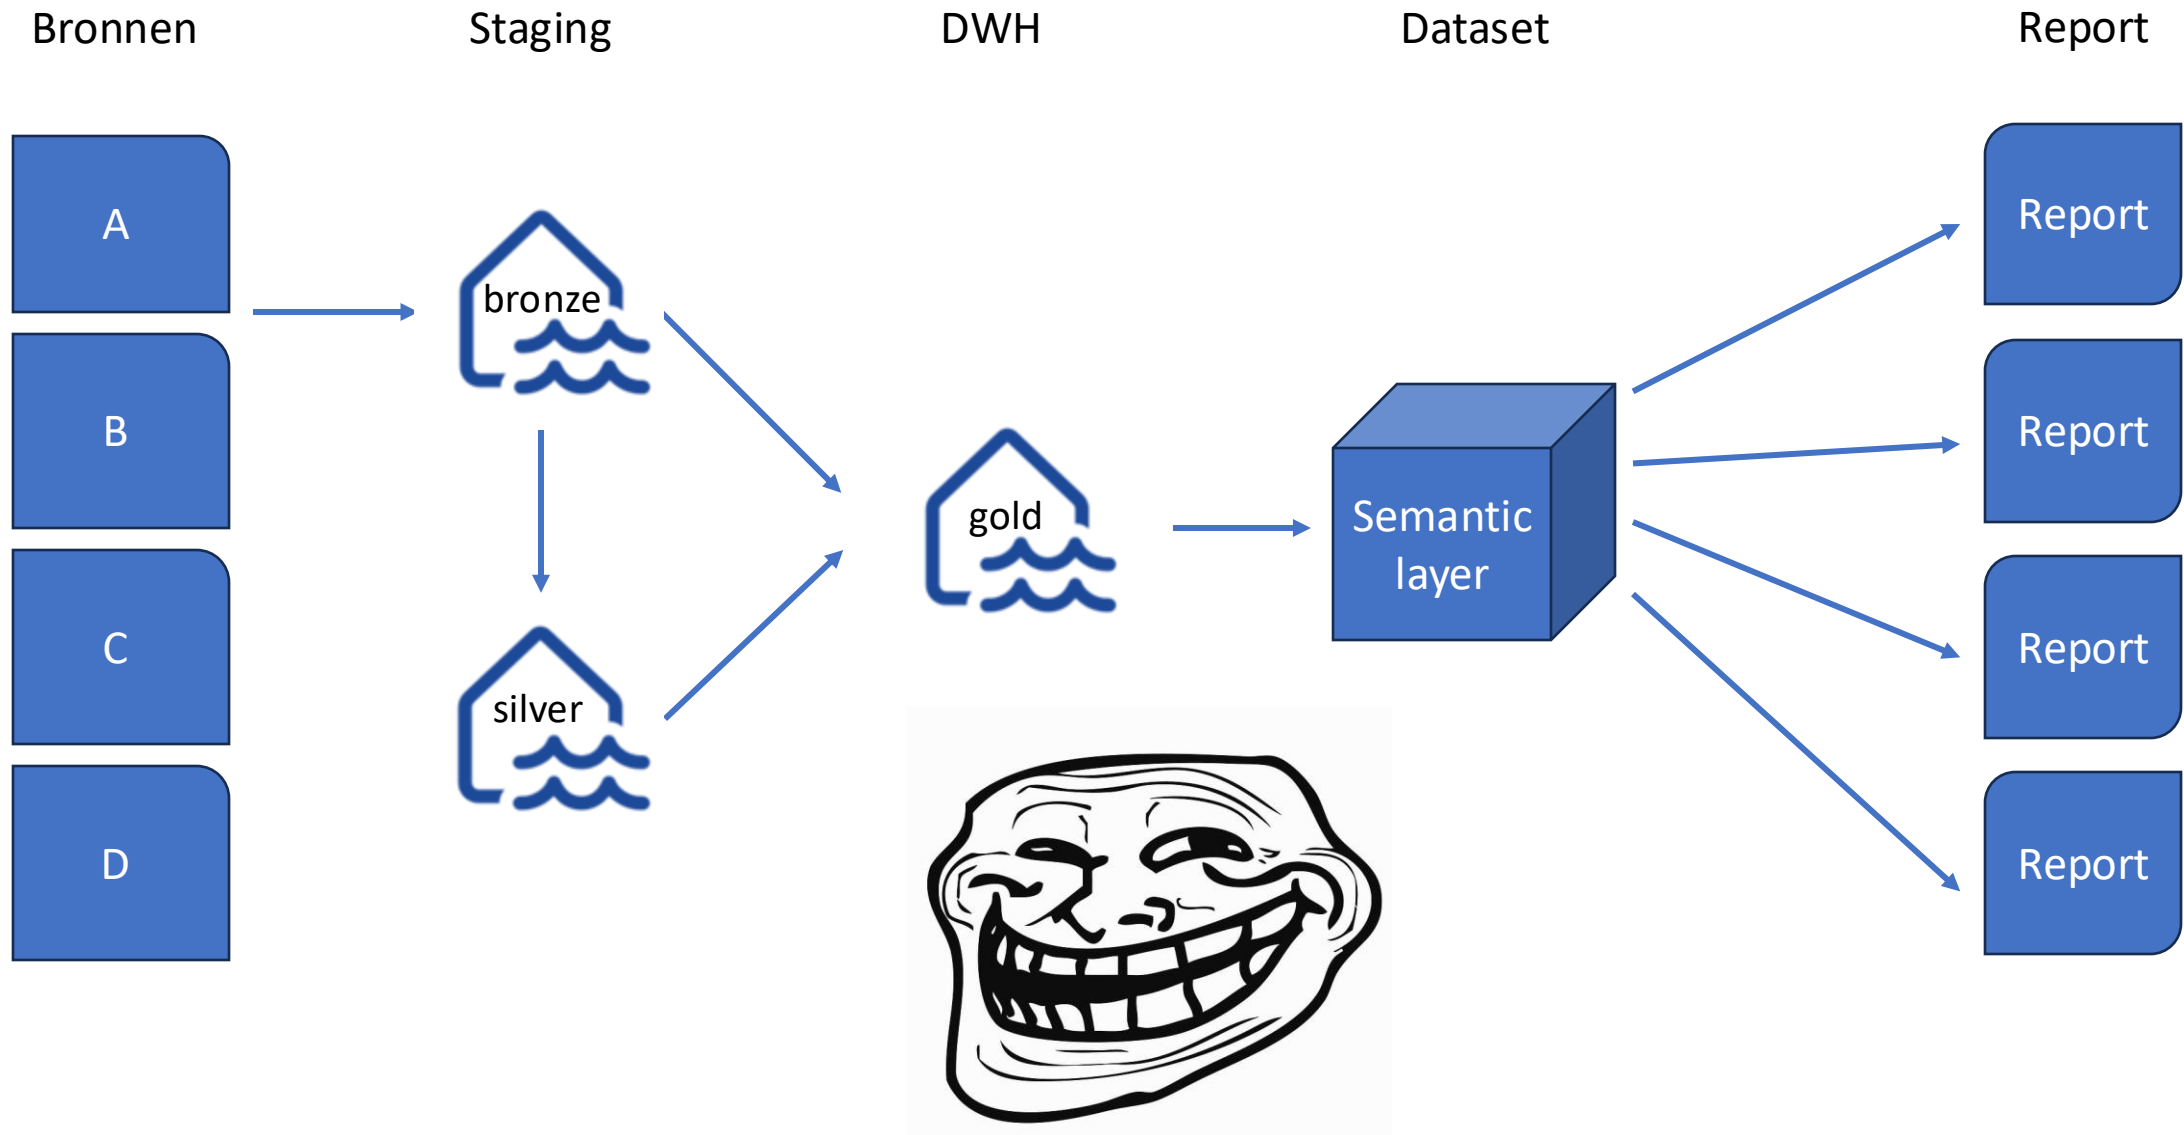
Design patterns

Topic 1/3

General architecture in the SQL world



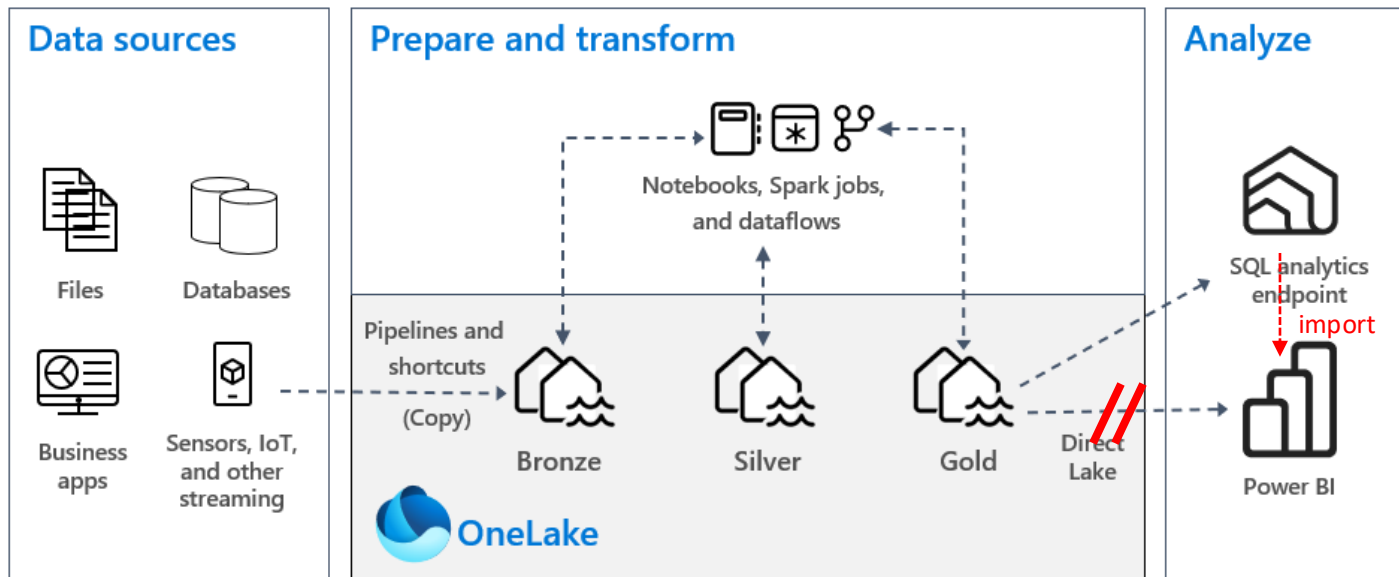
General architecture in the Lakehouse world



Medaillon architecture



- Not a Fabric-specific concept!
- Made popular by Databricks
- Describes three layers of data enrichment (bronze, silver, gold)



What we do:

- Bronze = raw data
- Silver = history, persistent stage, time travel, data types
- Gold = star schema

So... lakehouse vs database?



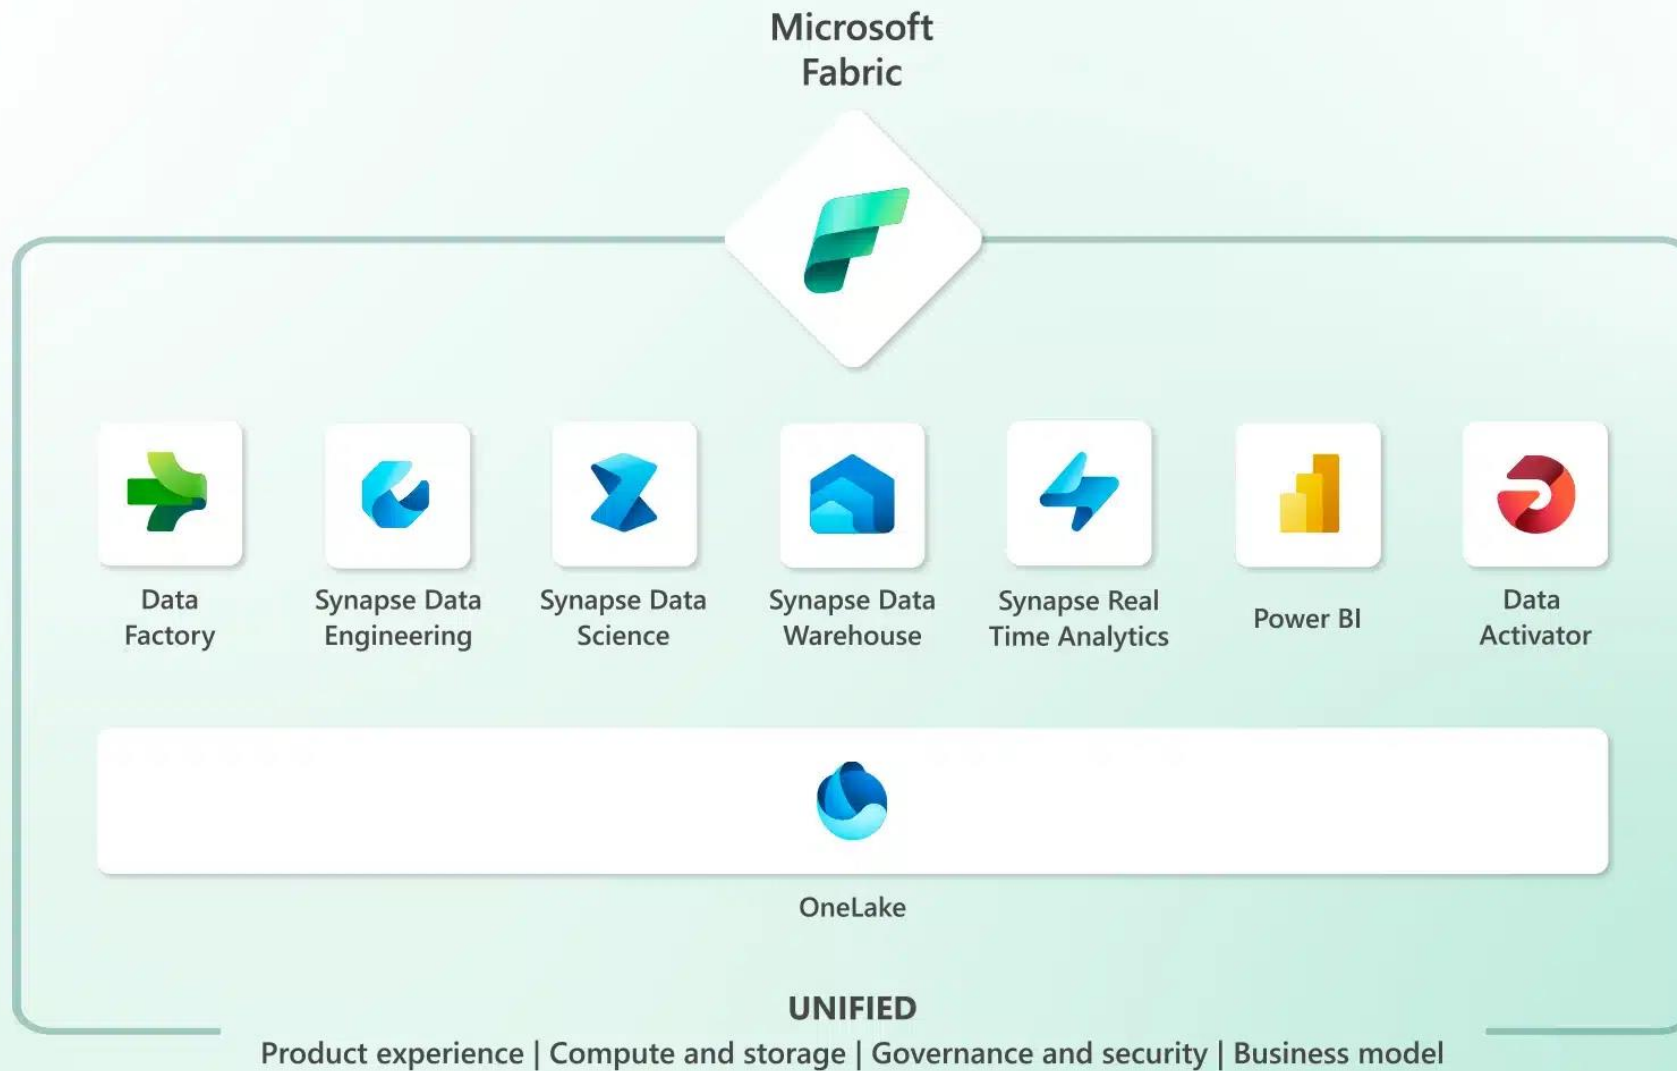
- We all know databases, right?
 - Structured data, with tables, columns, data types
- And data lakes:
 - Structured and unstructured data, such as csv, parquet, jpg, pdf, etc
- Lakehouse:
 - Best of both worlds?
 - Data lake storage, delta parquet tables with structures




Engines: SQL vs Spark

Topic 2/3

Remember this one?



Today's scope



Spark engine, OneLake storage (unstructured files, structured delta tables)



T-SQL engine, OneLake storage (structured delta tables), also, we don't use this one...
(perhaps only for configuration data)



Storage account, shortcuts, security

Fabric Warehouse



- Not used in Kimura – we only use Lakehouse
- Lakehouse = file (parquet / delta) storage with Spark engine
- Warehouse = file storage with Spark and T-SQL APIs
- T-SQL is amazing, but not full feature parity with SQL Server
- Plus, it's an abstraction on top of Spark
- Conclusion: learn to use (Py)Spark and go all-in on Lakehouses 😊
- With a bit of flexibility you can use SQL in your Lakehouses

Spark engine for data engineering



- First differences between SQL and Spark – processing languages
- In a database, it's SQL
- In Spark, we can choose whatever fits our requirements
 - Very popular: PySpark (Python for Spark) and SparkSQL

Spark engine: functional concepts



- Functional concepts don't change much
- For our Silver layer (previous: persistent stage), we now do Schema On Read instead of Schema On Write
 - This saves development time and lots of configurations

Spark engine: technical concepts



- Technical concepts do change a lot:
- Storage: from relational database to delta tables
- Data structure: from fixed schema tables to delta tables with transaction logs
- Language: T-SQL to Spark (variants)
- Style: transaction-based processing to distributed parallel processing

Spark languages



- PySpark (python) and SparkSQL (SQL) are a higher abstraction of native Spark
- You can work with DataFrames, and also with all the Python libraries you can think of
- Fabric has more than 100 available Python libraries, but you can also develop and install your own

PySpark notebooks for processing

- Kimura Data Framework: interaction between:
 - Kimura Python Library (data processing)
 - PySpark notebooks (data content, business rules, etc)

```
1 #Parameters
2 full_load = False
3 #if full_load = True then travel through all timestamps ignoring the timestamp param. This is not implemented yet.
4 timestamp = "20240229161122"#"20240206090005" #20240201055008
```

✓ - Apache Spark session ready in 3 min 13 sec 439 ms. Command executed in 2 sec 474 ms by admin-kimura on 5:23:05 PM, 2/29/24

```
1 #Imports
2 import kimura as k
3 import ast
4 from datetime import datetime
5 import json
6
7 #Base stuff
8 base_path = f"Files/Landing/Full"
```

✓ - Apache Spark session ready in 2 min 35 sec 458 ms. Command executed in 2 sec 64 ms by admin-kimura on 6:13:37 PM, 2/29/24

```
1 json_path = f"{base_path}/{timestamp}/metadata.json"
2 json_df = spark.read.json(json_path)
3 json_content = json_df.collect()[0].asDict()
4 table_list = json.loads(json_content['TableList'])
5 # display(table_list)
6
7 #Load data to Bronze layer delta tables
8 for table in table_list:
9     target_table = table['TargetTable']
10    source_system = table['SourceSystem']
11    file_type = table['FileType']
12    table_id = table['Id']
13    incremental_field = table['IncrementalField']
14    json_root_array = table['JSONRootArray']
15    table_key_list = ast.literal_eval(table['TableKey']) #
16
17    #Start logging
18    start_datetime = datetime.now()
19
20    #Execute
21    print(f"Starting with table: {target_table}")
22    count_rows = k.bronze_create_table(
23        timestamp = timestamp,
24        sourcesystem = source_system,
25        tablename = target_table,
26        tablekeys = table_key_list,
27        base_path = base_path,
28        file_type = file_type,
29        table_id = table_id,
30        incremental_field = incremental_field,
31        json_root_array = json_root_array,
32        spark_session = spark
33    )
34
35    #End logging
36    end_datetime = datetime.now()
37    delta = end_datetime - start_datetime
38    print(f"Table: {target_table}, rowcount: {count_rows}, duration in sec: {delta.total_seconds()}")
```

⊗ - Command executed in 1 min 29 sec 923 ms by admin-kimura on 5:24:36 PM, 2/29/24

Fabric notebooks



- Jupyter notebooks
- Support for Python, Scala, SparkSQL and R
 - Can be mixed!
 - Examples on the following slides

PySpark notebooks - PySpark



```
1 # load data using python, into a spark dataframe
2 spark_df = spark.read.format('delta').load('Tables/silver_dimDates')
3 display(spark_df)
```

✓ 17 sec - Command executed in 16 sec 775 ms by DWHDEV | Kimura on 1:56:40 PM, 4/20/24

> Spark jobs (2 of 2 succeeded) Resources Log

Table Chart | Showing rows 1 - 1000 ↓

	Date	123 Year	123 MonthNumber	ABC MonthName	123 DayOfMonth	123 DayOfWeek	ABC Day
1	2018-01-01	2018	1	January	1	2	Monday
2	2018-01-08	2018	1	January	8	2	Monday
3	2018-01-15	2018	1	January	15	2	Monday
4	2018-01-22	2018	1	January	22	2	Monday
5	2018-01-29	2018	1	January	29	2	Monday
6	2018-02-05	2018	2	February	5	2	Monday
7	2018-02-12	2018	2	February	12	2	Monday
8	2018-02-19	2018	2	February	19	2	Monday
9	2018-02-26	2018	2	February	26	2	Monday
10	2018-03-05	2018	3	March	5	2	Monday
11	2018-03-12	2018	3	March	12	2	Monday
12	2018-03-19	2018	3	March	19	2	Monday

PySpark notebooks – PySpark & SQL



```
1 # load data using python, into a spark dataframe, using a SQL string
2
3 query = "SELECT * FROM Silver.silver_dimDates LIMIT 10"
4
5 spark_df = spark.sql(query)
6 display(spark_df)
```

✓ 5 sec - Command executed in 5 sec 36 ms by DWHDEV | Kimura on 1:58:06 PM, 4/20/24

> Spark jobs (2 of 2 succeeded) Resources Log

Table Chart | Showing rows 1 - 10

	Date	123 Year	123 MonthNumber	ABC MonthName	123 DayOfMonth	123 DayOfWeek	ABC Day
1	2018-01-01	2018	1	January	1	2	Monday
2	2018-01-08	2018	1	January	8	2	Monday
3	2018-01-15	2018	1	January	15	2	Monday
4	2018-01-22	2018	1	January	22	2	Monday
5	2018-01-29	2018	1	January	29	2	Monday
6	2018-02-05	2018	2	February	5	2	Monday
7	2018-02-12	2018	2	February	12	2	Monday
8	2018-02-19	2018	2	February	19	2	Monday
9	2018-02-26	2018	2	February	26	2	Monday

PySpark notebooks – SparkSQL



```
1  %%sql
2  -- query data using SQL
3
4  SELECT
5      t.*
6  FROM Silver.silver_dimDates t
7  LIMIT 10
```

✓ 2 sec - Command executed in 1 sec 573 ms by DWHDEV | Kimura on 1:59:44 PM, 4/20/24

> Spark jobs (2 of 2 succeeded) Resources

Table Chart

	Date	123 Year	123 MonthNumber	ABC MonthName	123 DayOfMonth	123 DayOfWeek	ABC Day
1	2018-01-01	2018	1	January	1	2	Monday
2	2018-01-08	2018	1	January	8	2	Monday
3	2018-01-15	2018	1	January	15	2	Monday
4	2018-01-22	2018	1	January	22	2	Monday
5	2018-01-29	2018	1	January	29	2	Monday
6	2018-02-05	2018	2	February	5	2	Monday
7	2018-02-12	2018	2	February	12	2	Monday
8	2018-02-19	2018	2	February	19	2	Monday
9	2018-02-26	2018	2	February	26	2	Monday
10	2018-03-05	2018	3	March	5	3	Monday

PySpark notebooks – SparkSQL -> PySpark

```
1  %%sql
2  CREATE OR REPLACE TEMPORARY VIEW my_temp_view AS
3  SELECT
4      t.*
5  FROM Silver.silver_dimDates t
6  LIMIT 10
```


✓ 3 sec - Command executed in 2 sec 405 ms by DWHDEV | Kimura on 2:18:05 PM, 4/20/24





>  Log



No data available

```
1  spark_df_from_view = spark.sql("select * from my_temp_view")
2  display(spark_df_from_view)
```

✓ 2 sec - Command executed in 1 sec 588 ms by DWHDEV | Kimura on 2:18:32 PM, 4/20/24

>  Spark jobs (2 of 2 succeeded)  Resources  Log

 Table  Chart | Showing rows 1 - 10  

	 Date	123 Year	123 MonthNumber	ABC MonthName	123 DayOfMonth	123 DayOfWeek
1	2018-01-01	2018	1	January	1	2
2	2018-01-08	2018	1	January	8	2
3	2018-01-15	2018	1	January	15	2
4	2018-01-22	2018	1	January	22	2
5	2018-01-29	2018	1	January	29	2
6	2018-02-05	2018	2	February	5	2
7	2018-02-12	2018	2	February	12	2
8	2018-02-19	2018	2	February	19	2
9	2018-02-26	2018	2	February	26	2
10	2018-03-05	2018	3	March	5	1

PySpark notebooks – PySpark -> SparkSQL

```
1 # load data using python, into a spark dataframe
2 spark_df = spark.read.format('delta').load('Tables/silver_dimDates')
3 spark_df.createOrReplaceTempView('my_temp_view')
```

✓ 1 sec - Command executed in 838 ms by DWHDEV | Kimura on 2:21:11 PM, 4/20/24



>  Log

```
1 %%sql
2 SELECT
3     t.*
4 FROM my_temp_view t
5 LIMIT 10
```

✓ 2 sec - Command executed in 1 sec 489 ms by DWHDEV | Kimura on 2:21:16 PM, 4/20/24

>  Spark jobs (2 of 2 succeeded)  Resources  Log

 Table  Chart | 

	 Date	123 Year	123 MonthNumber	ABC MonthName	123 DayOfMonth	123 DayOfWeek	ABC Day
1	2018-01-01	2018	1	January	1	2	Monday
2	2018-01-08	2018	1	January	8	2	Monday
3	2018-01-15	2018	1	January	15	2	Monday
4	2018-01-22	2018	1	January	22	2	Monday
5	2018-01-29	2018	1	January	29	2	Monday
6	2018-02-05	2018	2	February	5	2	Monday
7	2018-02-12	2018	2	February	12	2	Monday
8	2018-02-19	2018	2	February	19	2	Monday



Code reusability

Topic 3/3

Code reusability



- Manual work vs Code generation vs Code reusability
- Manual work: NOPE
- Code generation: cool, done that a lot, but it's not needed anymore
- Code reusability: write code once, execute many times
- Examples on the next slides

Lazy is good



- In Fabric, we try to be lazy data engineers
- The more work we do manually, the less customers we can help and the more errors we make
- The general rule is; everything you do twice, you put into our Kimura Data Framework
- DRY: Don't Repeat Yourself

Generated code

```
SELECT
    ISNULL([t].id, '') AS [Id]
    , ISNULL([t].internalName, '') AS [InternalName]
    , ISNULL([t].externalName, '') AS [ExternalName]
    , ISNULL([t].carrierId, 0) AS [CarrierId]
    , ISNULL([t].country, '') AS [Country]
    , ISNULL([t].batchNumber, 0) AS [BatchNumber]
    , ISNULL([t].deletedAt, '1900-01-01') AS [DeletedAt]
    , ISNULL(getDate(), '1900-01-01') AS [Sys_FirstLoadTimeStamp]
    , ISNULL(getDate(), '1900-01-01') AS [Sys_LastUpdateTimeStamp]
    , ISNULL('I', '') AS [Sys_Operation]
    , ISNULL('1900-01-01', '1900-01-01') AS [Sys_ValidFrom]
    , ISNULL('2099-12-31', '1900-01-01') AS [Sys_ValidTo]
    , ISNULL(1, 1) AS [Sys_CurrentRow]
INTO #temp
FROM [pers].[Homerr-production_Depots] as t
where t.Sys_ValidFrom <= @date and t.Sys_ValidTo > @date

create index ix_temp on #temp ([Id])

MERGE INTO [dwh].[dimDepots] AS target
USING #temp AS source
ON target.[Id] = source.[Id]
```

```
42 WHEN MATCHED AND (
43     target.[Id] <> source.[Id]
44 OR   target.[InternalName] <> source.[InternalName]
45 OR   target.[ExternalName] <> source.[ExternalName]
46 OR   target.[CarrierId] <> source.[CarrierId]
47 OR   target.[Country] <> source.[Country]
48 OR   target.[BatchNumber] <> source.[BatchNumber]
49 OR   target.[DeletedAt] <> source.[DeletedAt]
50 ) THEN
51     UPDATE SET
52         target.[Id] = source.[Id]
53         , target.[InternalName] = source.[InternalName]
54         , target.[ExternalName] = source.[ExternalName]
55         , target.[CarrierId] = source.[CarrierId]
56         , target.[Country] = source.[Country]
57         , target.[BatchNumber] = source.[BatchNumber]
58         , target.[DeletedAt] = source.[DeletedAt]
59         , target.[Sys_LastUpdateTimeStamp] = source.[Sys_LastUpdateTimeStamp]
60         , target.[Sys_Operation] = 'U'
61         , target.[Sys_StageDB] = source.[Sys_StageDB]
62
63     WHEN NOT MATCHED THEN
64         INSERT (
65             [Id]
66             , [InternalName]
67             , [ExternalName]
68             , [CarrierId]
69             , [Country]
70             , [BatchNumber]
71             , [DeletedAt]
72             , [Sys_FirstLoadTimeStamp]
73             , [Sys_LastUpdateTimeStamp]
74             , [Sys_Operation]
75             , [Sys_StageDB]
76             , [Sys_ValidFrom]
77             , [Sys_ValidTo]
78             , [Sys_CurrentRow]
79         )
80         VALUES(
81             source.[Id]
82             , source.[InternalName]
83             , source.[ExternalName]
84             , source.[CarrierId]
85             , source.[Country]
86             , source.[BatchNumber]
87             , source.[DeletedAt]
88             , source.[Sys_FirstLoadTimeStamp]
89             , source.[Sys_LastUpdateTimeStamp]
90             , 'I'
91             , source.[Sys_StageDB]
92             , source.[Sys_ValidFrom]
93             , source.[Sys_ValidTo]
94             , source.[Sys_CurrentRow]
95         );
```



Reusable functions...



```
9 #Generic path names for Silver tables
10 def silver_target_path(silver_table: str) -> str:
11     path = f"Tables/silver_{silver_table}"
12     return path
13
14 #Create Sys fields
15 def add_sys_fields(df: sdf, business_key_columns: list) -> sdf:
16     df_with_sys_fields = \
17         df.withColumn("Sys_ID", sha2(concat_ws("~", *business_key_columns), 256)) \
18         .withColumn("Sys_ValidFrom", to_timestamp(lit("1900-01-01 00:00:00.000"), "yyyy-MM-dd HH:mm:ss.SSS")) \
19         .withColumn("Sys_ValidTo", to_timestamp(lit("2999-12-31 00:00:00.000"), "yyyy-MM-dd HH:mm:ss.SSS")) \
20         .withColumn("Sys_IsActive", lit(1).cast(BooleanType())) \
21         .withColumn("Sys_LoadDatetime", current_timestamp())
22     return df_with_sys_fields
23
24 #Upsert into delta table logic
25 def silver_create_table(silver_table: str, df_mapping: sdf, business_key_columns: list, spark_session: SparkSession) -> str:
26     """
27     This function performs an upsert into a delta table in Silver. If the table doesn't exist, it will be created.
28     """
29
30     result = "" #hiermee kunnen we bijvoorbeeld logging output geven
31
32     df_to_merge = add_sys_fields(df_mapping, business_key_columns)
33
34     target_path = silver_target_path(silver_table)
35     exclude_from_update = ['Sys_ID', 'Sys_ValidFrom', 'Sys_ValidTo', 'Sys_IsActive', 'Sys_LoadDatetime']
36     update_cols = {col: f"s.`{col}`" for col in df_to_merge.columns if col not in exclude_from_update}
37
38     if DeltaTable.isDeltaTable(spark_session, target_path):
39         silver_table = DeltaTable.forPath(spark_session, target_path)
40
41         #Upsert dataframe to delta
42         update_condition = " AND ".join(f"t.`{key_col}` = s.`{key_col}`" for key_col in business_key_columns)
43         silver_table.alias("t").merge(
44             source = df_to_merge.alias("s"),
45             condition = update_condition
46         ).whenMatchedUpdate(set=update_cols).whenNotMatchedInsertAll().execute()
47     else:
48         #if silver table doesn't exist yet
49         df_to_merge.write.format('delta').mode('overwrite').save(target_path)
50
51     return result
```

...make for super easy notebooks



```
1  import kimura as k
2
3  #Mapping logic
4  silver_table = "dimHandlingTypes"
5  business_key_columns = ['HandlingTypeId']
6
7  query = """
8  select DISTINCT
9      t.CARGOHANDLINGTYPE as HandlingTypeId
10     ,e.MEMBERNAME as Description
11  from Bronze.bronze_BYOD_FlxUS_CargoTransHandlingStaging t
12  left join Bronze.bronze_BYOD_FlxUS_EnumValueTableStaging e
13      on t.CARGOHANDLINGTYPE = e.VALUE
14      and e.ENUMNAME = 'FlxUs_CargoHandlingType'
15  """
16
17  df_mapping = spark.sql(query)
18
19  k.silver_create_table(
20      silver_table = silver_table,
21      df_mapping = df_mapping,
22      business_key_columns = business_key_columns,
23      spark_session = spark
24  )
```



Your first lakehouse

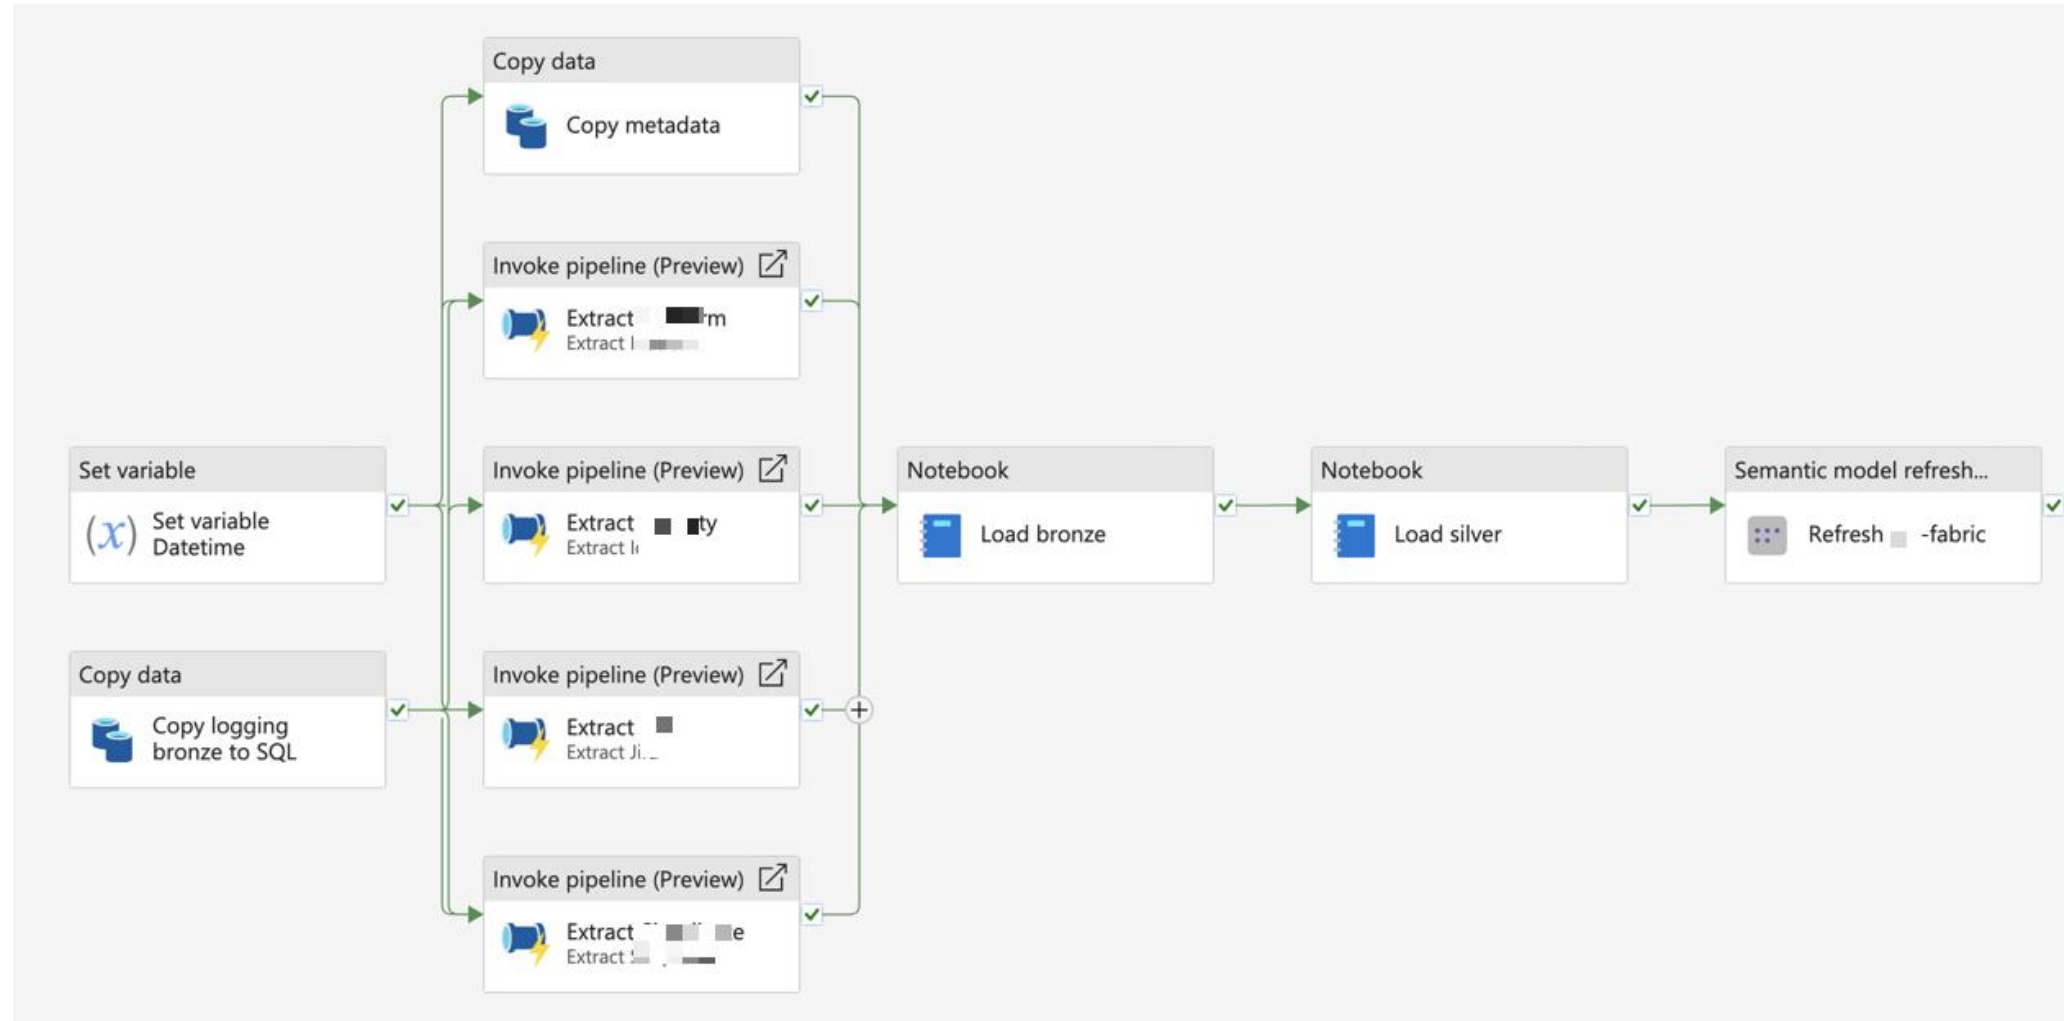
End-to-end overview

Rulebook



- Use Data Factory Pipelines for ingest and orchestration
 - Parametrised pipelines using Azure SQL (or Fabric Warehouse?) config database
- Use Lakehouse for storage (files in bronze, delta in silver + gold)
- PySpark and SparkSQL in notebooks for processing
 - Reusable code with custom Python library
- Power BI semantic model using Import Mode and SQL endpoint on Gold Lakehouse

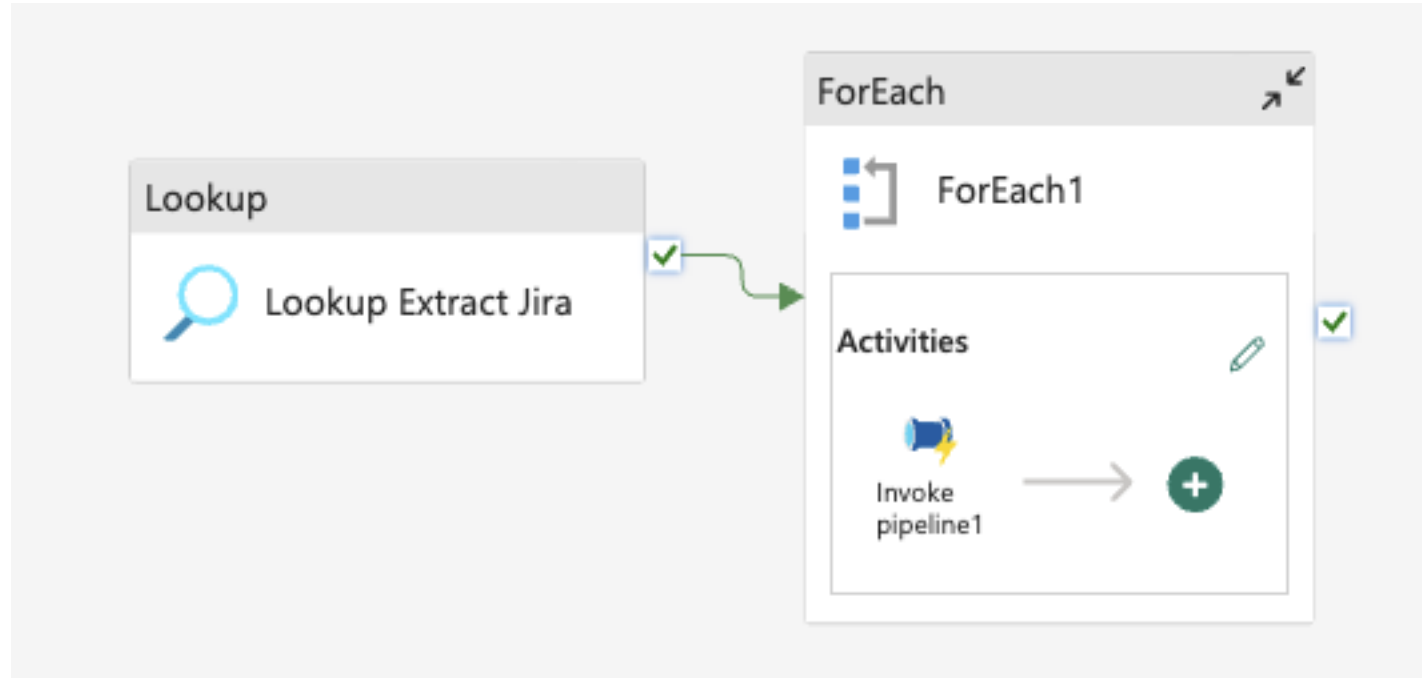
Main pipeline (simplified)



Ingest / extract example



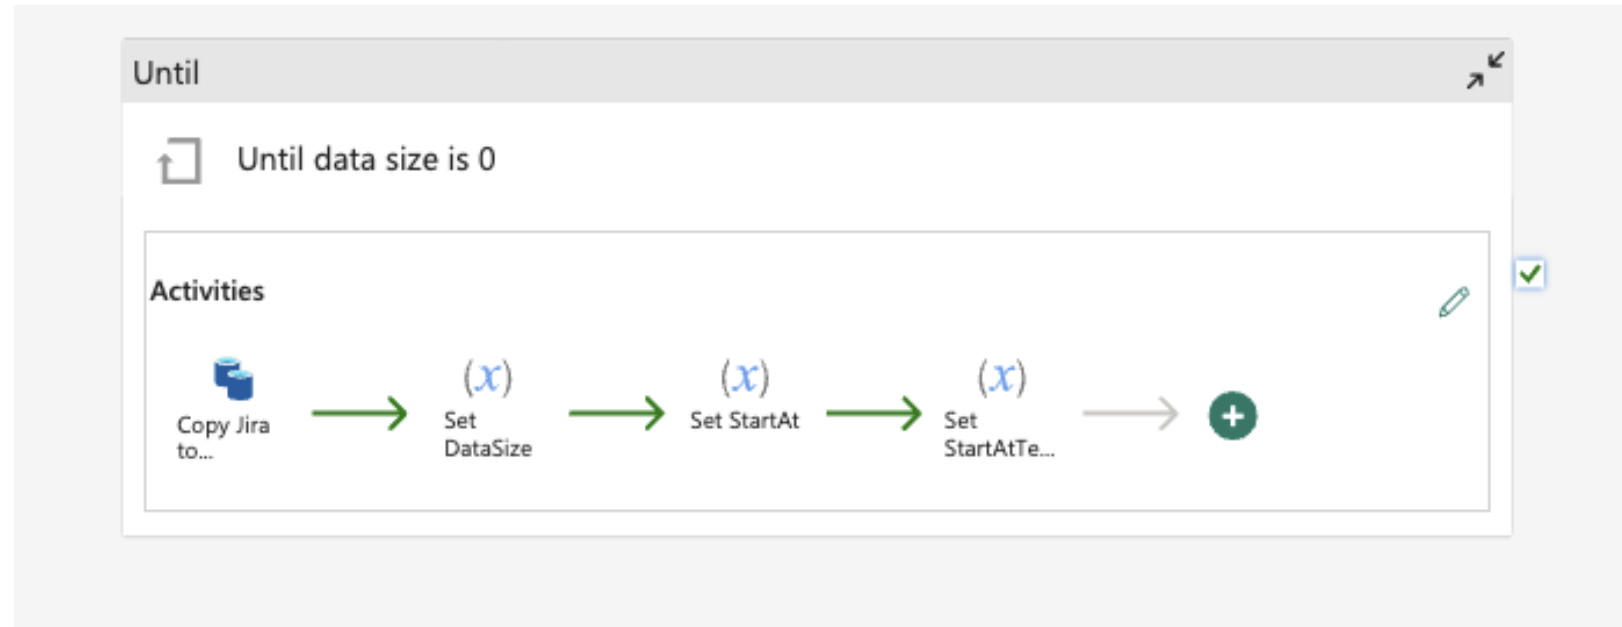
- Pipeline + data lake folder structure



Ingest / extract example



- Pipeline + data lake folder structure





Ingest / extract example

- Pipeline + data lake folder structure

Pipeline expression builder



Add dynamic content below using any combination of [expressions](#), [functions](#) and [system variables](#).

```
@concat(  
  'Landing/',  
  '/Full/',  
  pipeline().parameters.Datetime,  
  '/',  
  pipeline().parameters.SourceSystem,  
  '/',  
  pipeline().parameters.TargetTable  
)
```

[Clear contents](#)

Notebooks



- Starting with a very light weight setup, from your browser:

The screenshot displays the Microsoft Power BI web interface in a browser window. The address bar shows 'app.powerbi.com' and the page title is 'silver_load - Power BI'. The interface includes a left-hand navigation pane with a sidebar menu containing icons for Start, Maken, Bladeren, OneLake-gegevens..., Apps, Metrieken, Monitor, Informatie, Realtime-hub, Werkruimte, DWH DEV, and silver_load. The main content area is titled 'silver_load | Opgeslagen' and features a search bar. Below the search bar, there are tabs for 'Start', 'Bewerken', 'Uitvoeren', and 'Weergeven'. The 'Start' tab is active, showing a list of 'Lakehouses' and 'Silver' data sources. A warning message at the top of the main area states: 'Andere personen in uw organisatie hebben mogelijk toegang tot notebooks en Spark-taakdefinities in deze werkrumte. Controleer dit item zorgvuldig voordat u het uitvoert.' The main workspace contains three code blocks. The first block is titled '#Parameters' and shows a single line of code. The second block is titled '#Imports' and shows code for importing libraries like 'kimura', 'pandas', and 'datetime'. The third block is titled '#Mapping logic' and shows code for setting up data ranges and mappings. The code is written in Python and uses Spark syntax. The interface also includes a bottom status bar with 'PySpark (Python)' and 'PySpark (Python)' labels.

Notebooks



- Starting with a very light weight setup, from your browser:
- You can worry about Visual Studio Code later 😊

Notebook Bronze -> Silver



```
1 json_path = f"{base_path}/{timestamp}/metadata.json"
2 json_df = spark.read.json(json_path)
3 json_content = json_df.collect()[0].asDict()
4 table_list = json.loads(json_content['TableList'])
5 # display(table_list)
6
7 #Load data to Bronze layer delta tables
8 for table in table_list:
9     target_table = table['TargetTable']
10    source_system = table['SourceSystem']
11    file_type = table['FileType']
12    table_id = table['Id']
13    incremental_field = table['IncrementalField']
14    json_root_array = table['JSONRootArray']
15    table_key_list = ast.literal_eval(table['TableKey']) #
16
17    #Start logging
18    start_datetime = datetime.now()
19
20    #Execute
21    print(f"Starting with table: {target_table}")
22    count_rows = k.bronze_create_table(
23        timestamp = timestamp,
24        sourcesystem = source_system,
25        tablename = target_table,
26        tablekeys = table_key_list,
27        base_path = base_path,
28        file_type = file_type,
29        table_id = table_id,
30        incremental_field = incremental_field,
31        json_root_array = json_root_array,
32        spark_session = spark
33    )
34
35    #End logging
36    end_datetime = datetime.now()
37    delta = end_datetime - start_datetime
38    print(f"Table: {target_table}, rowcount: {count_rows}, duration in sec: {delta.total_seconds()}")
```

```
1 #Parameters
2 full_load = False
3 #if full_load = True then travel thr
4 timestamp = "20240223131249"#"202402
```

✓ - Apache Spark session ready in 3 min 30 sec 150 ms. Comman

```
1 #Imports
2 import kimura as k
3 import ast
4 from datetime import datetime
5 import json
6
7 #Base stuff
8 base_path = f"Files/Landing/Full"
```

✓ - Command executed in 270 ms by DWHDEV | Kimura on 3:57:4

Notebook Silver -> Gold



```
1  #Mapping logic
2  silver_table = "dimProjects"
3  business_key_columns = ['ProjectId']
4
5  query = """
6  select
7      t.`data.id` as ProjectId
8      ,t.`data.name` as ProjectName
9      ,regexp_replace(t.`data.project_status.label`,`tab_`,`') as Status
10     ,t.`data.organization_id` as OrganisationId
11     ,t.`data.start_date` as StartDate
12     ,t.`data.end_date` as EndDate
13     ,c.Sys_ID as FK_dimCustomers_Sys_ID
14  from Bronze.bronze_Simplicate_ProjectsProject t
15  left join Silver.silver_dimCustomers c
16      on t.`data.organization.id` = c.SimplicateId
17  where 1=1
18  """
19
20  df_mapping = spark.sql(query)
21  # display(df_mapping)
22  output = k.silver_create_table(silver_table = silver_table, df_mapping = df_map
```

Recap



- Fabric Lakehouses are very powerful
- Spark (especially SparkSQL) is not hard to learn for SQL Engineers
- Code reusability is much easier than in SQL
- So... start using it tomorrow?



That's it!

Questions?

Kimura Data Intelligence B.V.

“Wij zoeken dringend een lead data consultant. Ben jij diegene?”

Fonteinkruid 6b
3931 WX, Woudenberg
The Netherlands

www.kimura.nl
info@kimura.nl

